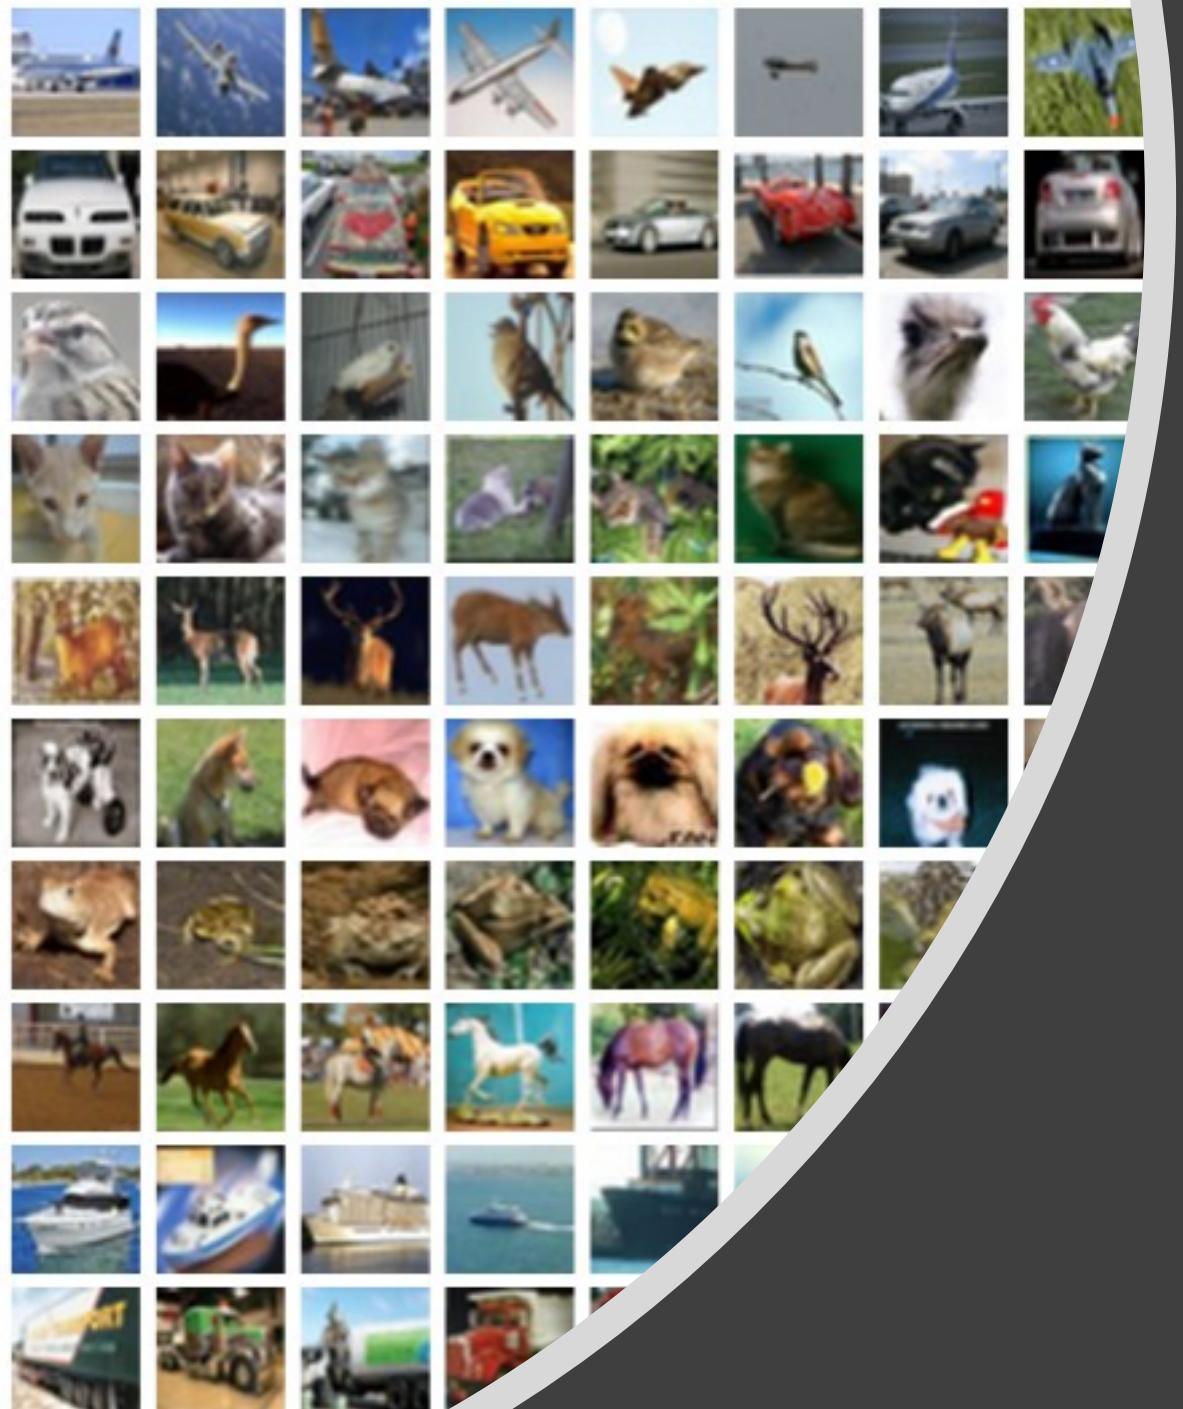


le



# Model Fitting

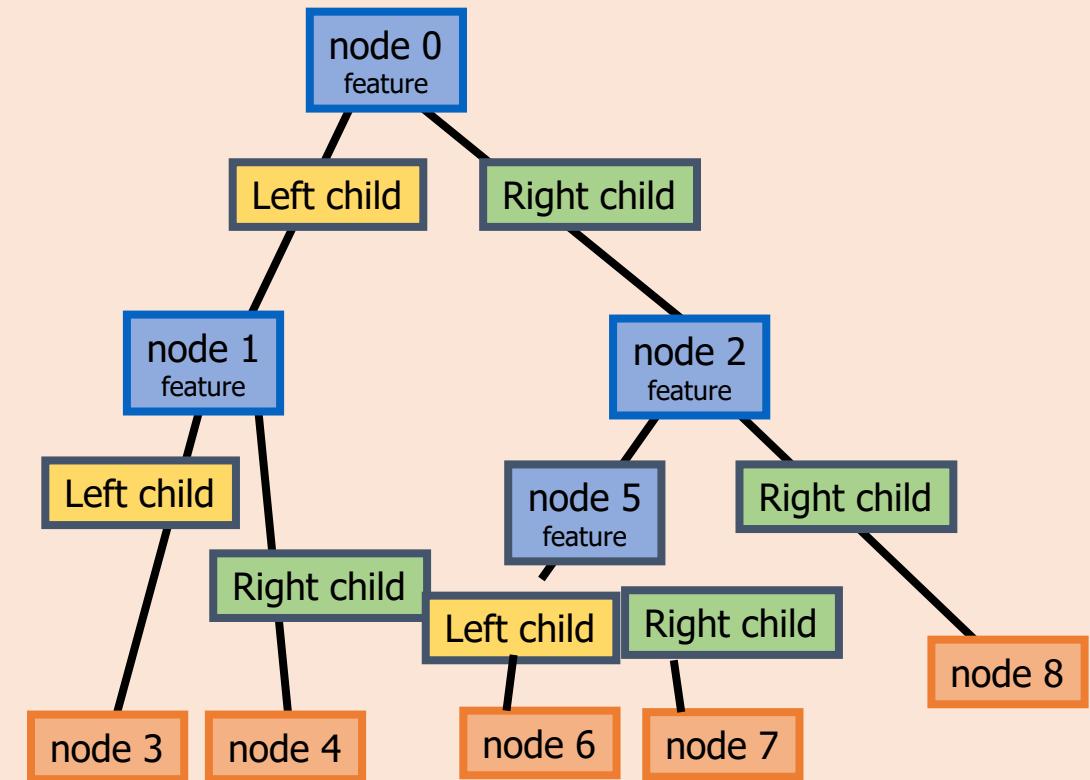
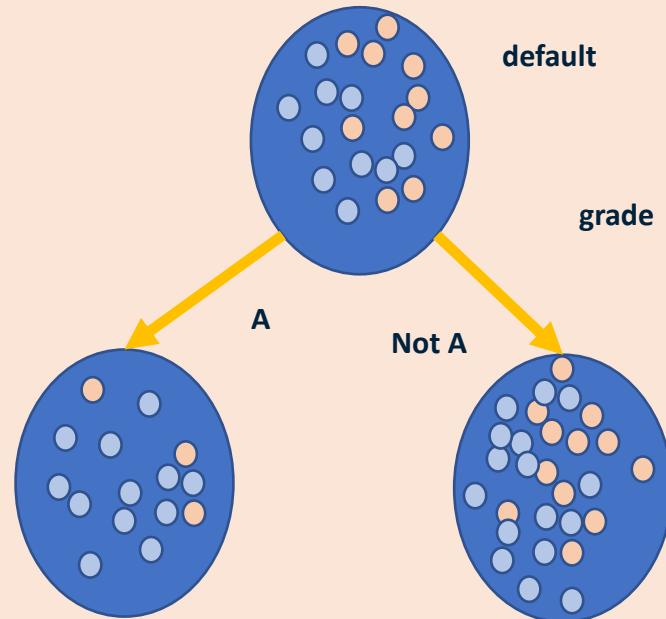
## Classifying Cats



- Decision Tree recap
- Geometry
- Linear models
- Linear Classifiers
- Logistic regression
- Implementation

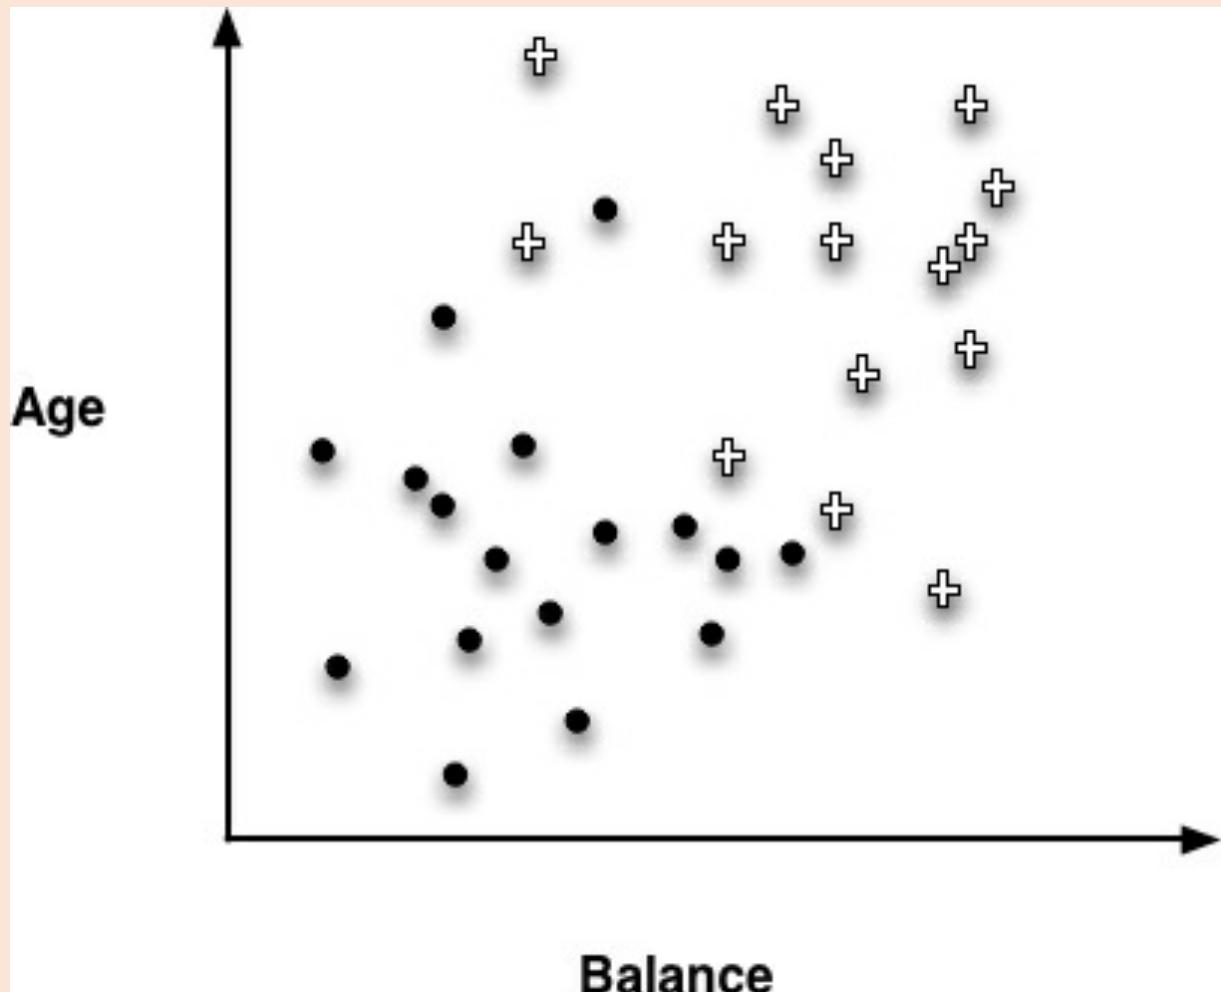
# Decision Tree

## Where we left off



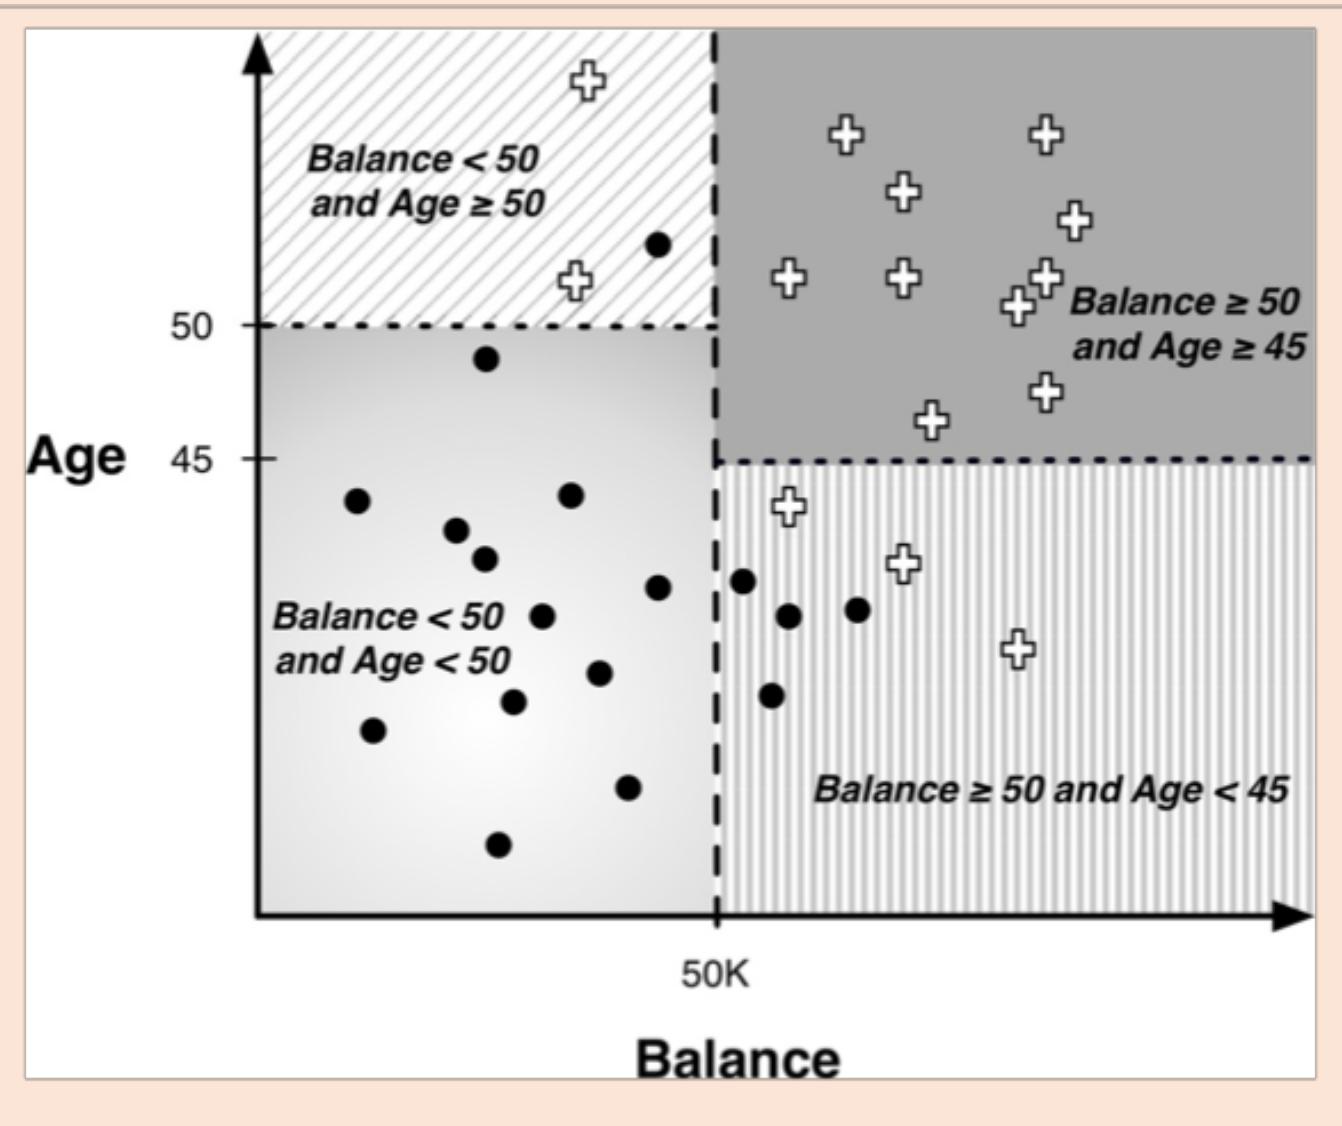
## Geometric Interpretation

---



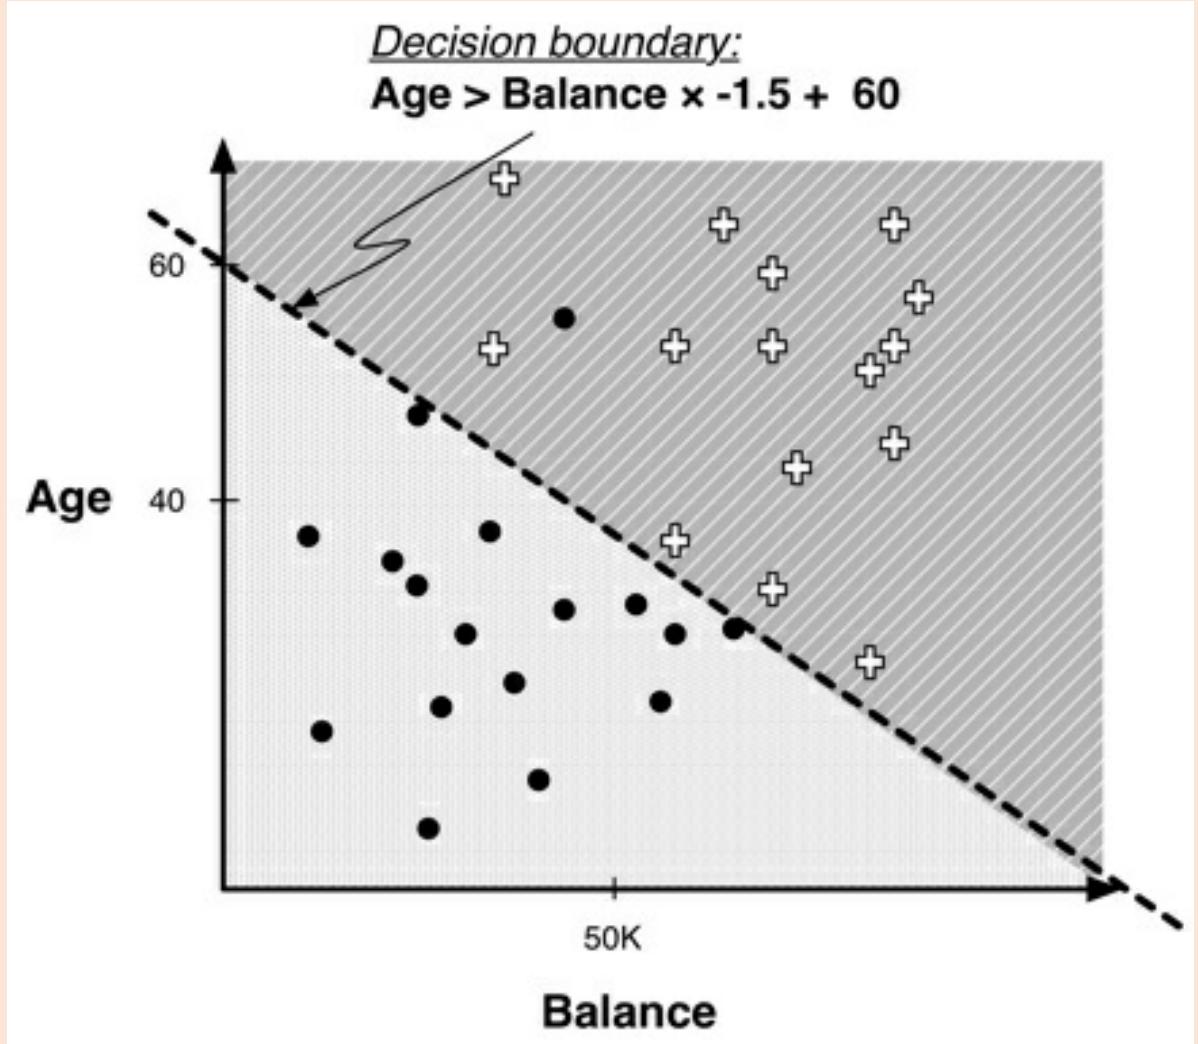
## Geometric Interpretation

---



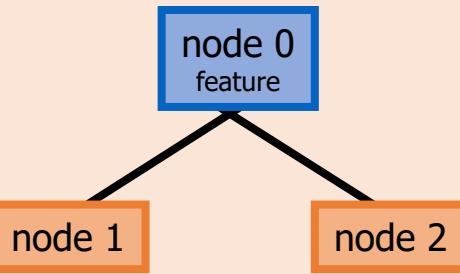
## Geometric Interpretation

### Linear Regression

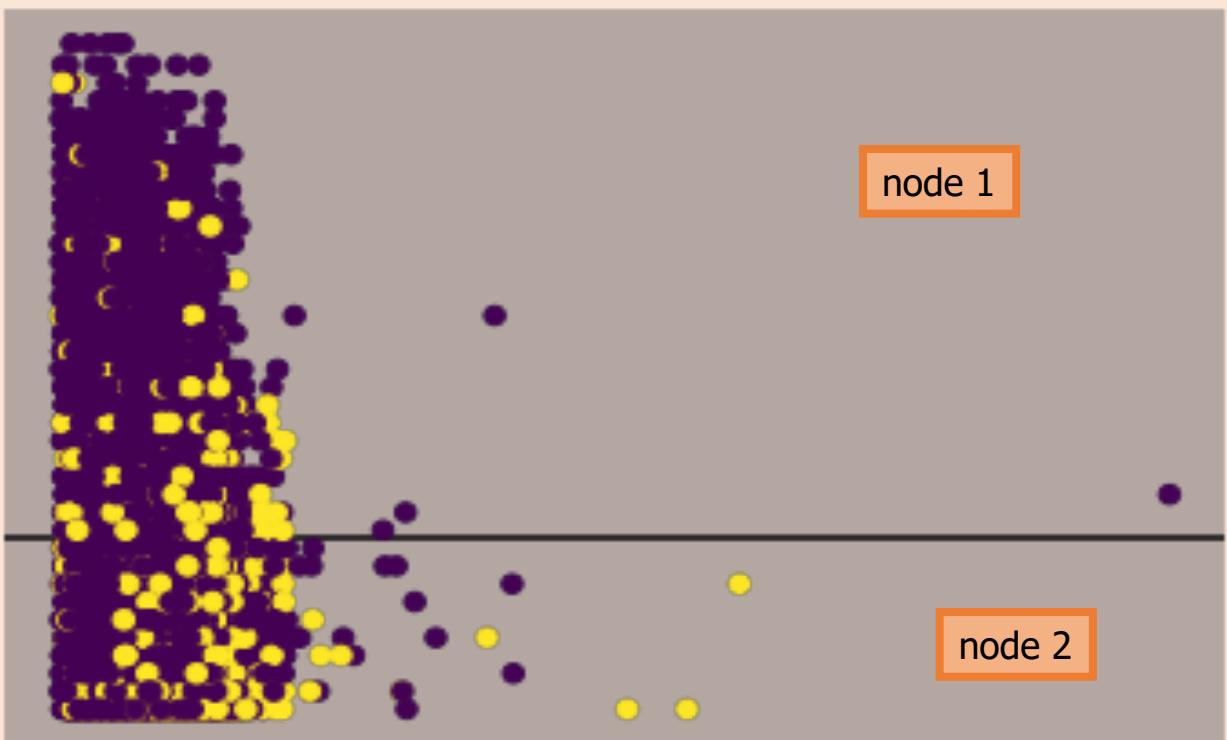


# Geometric Interpretation

---

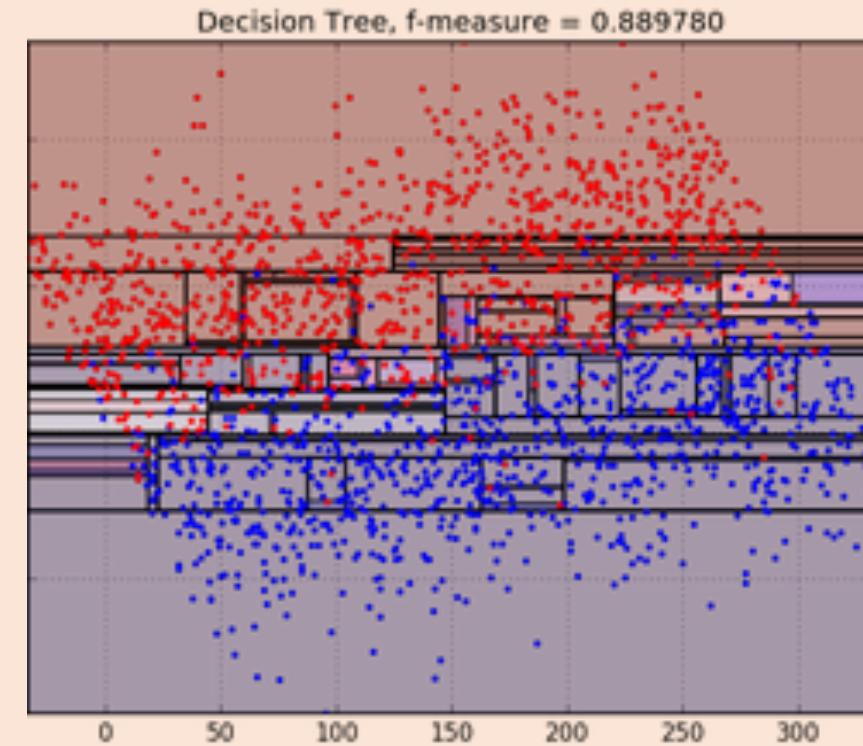
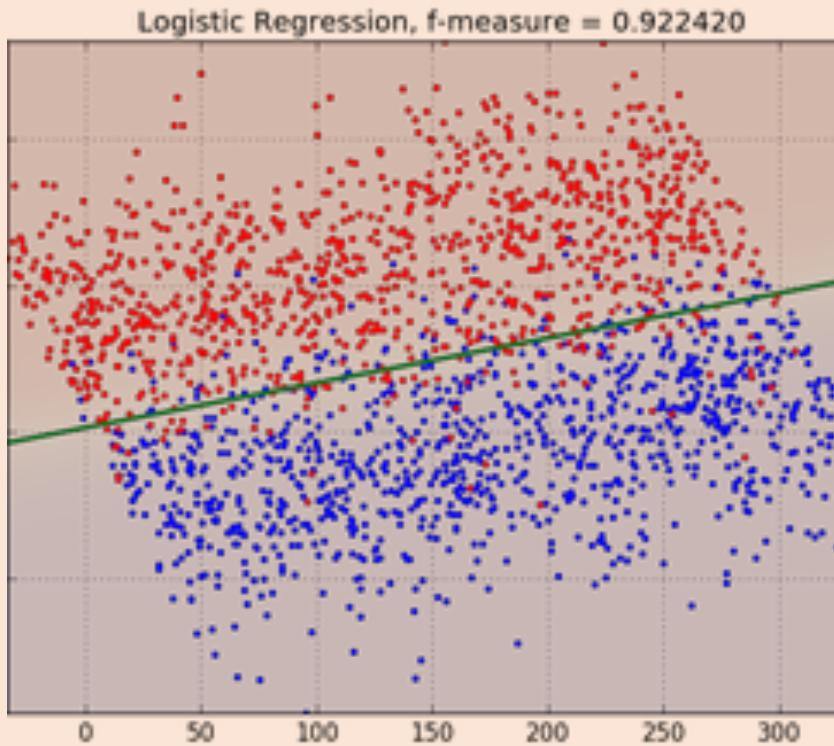


# leaf nodes = 2



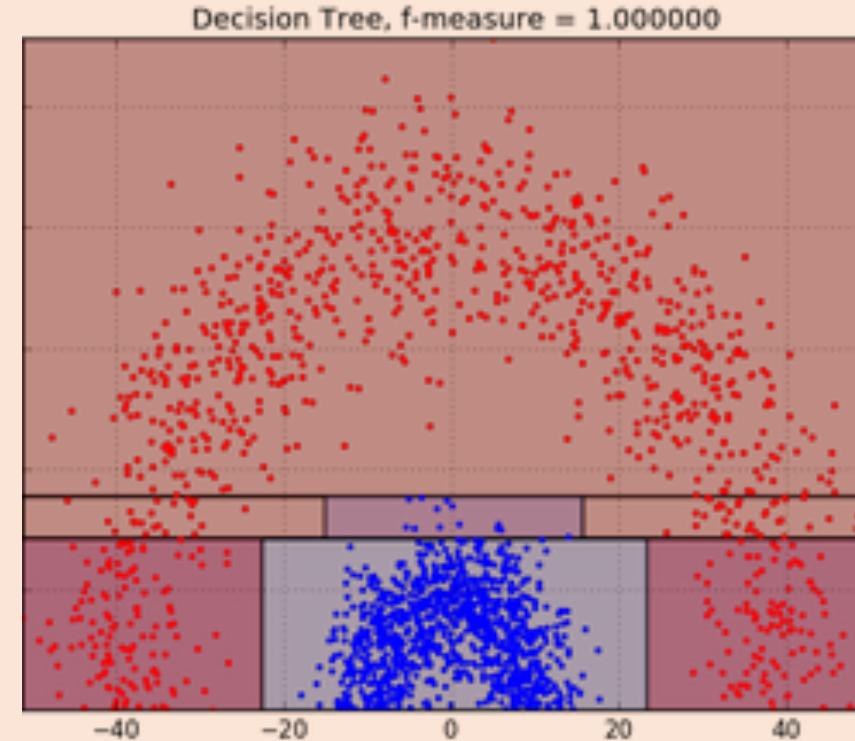
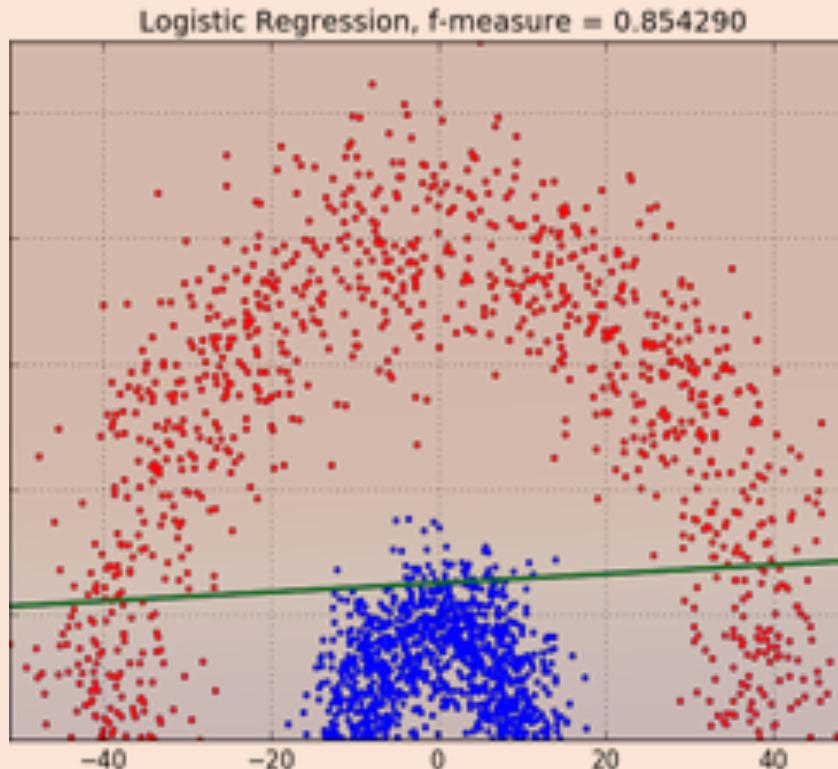
# Geometric Interpretation

---

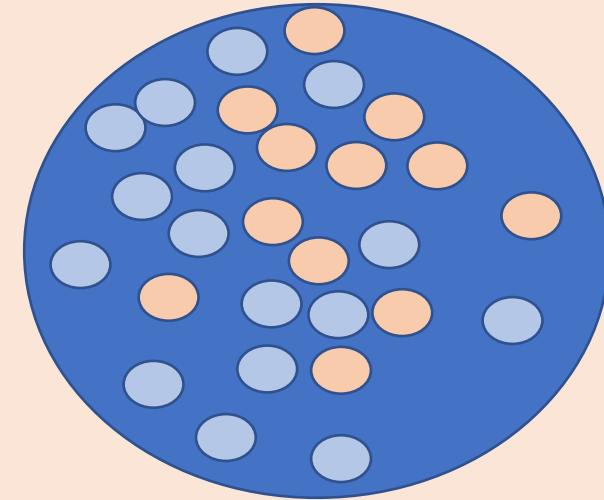


# Geometric Interpretation

---



# Parametric and Non parametric

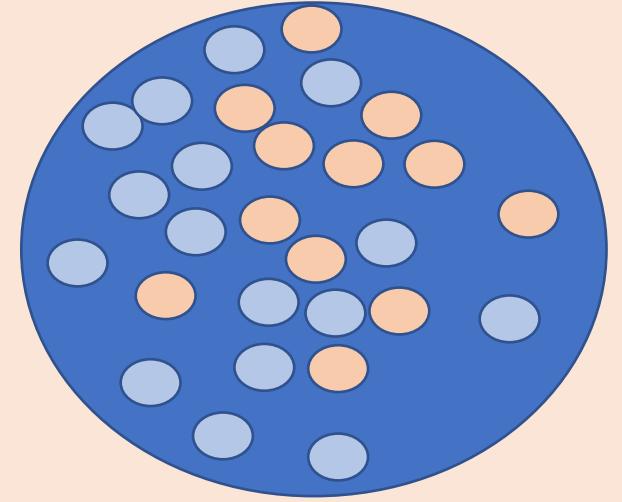


In all cases our objective is to find a model/forecast/function that best fits the data in the sense of minimizing loss

- A **non parametric** form does not make any assumptions on the functional form
- A **parametric** form assumes a functional form and identifies the parameters that minimize the loss within the category

# LOSS

1. Define a **loss function** that quantifies our unhappiness with the scores across the training data.
2. Come up with a way of efficiently finding the parameters that minimize the loss function. (**optimization**)



Let  $\hat{y}_1 \dots \hat{y}_k$  be as before and  $p$  the estimated parameter

[ $L_2$  - loss]: 
$$L(\hat{y}) = \sum_i (\hat{y}_i - p)^2$$

[ $L_1$  - loss]: 
$$L(\hat{y}) = \sum_i |\hat{y}_i - p|$$

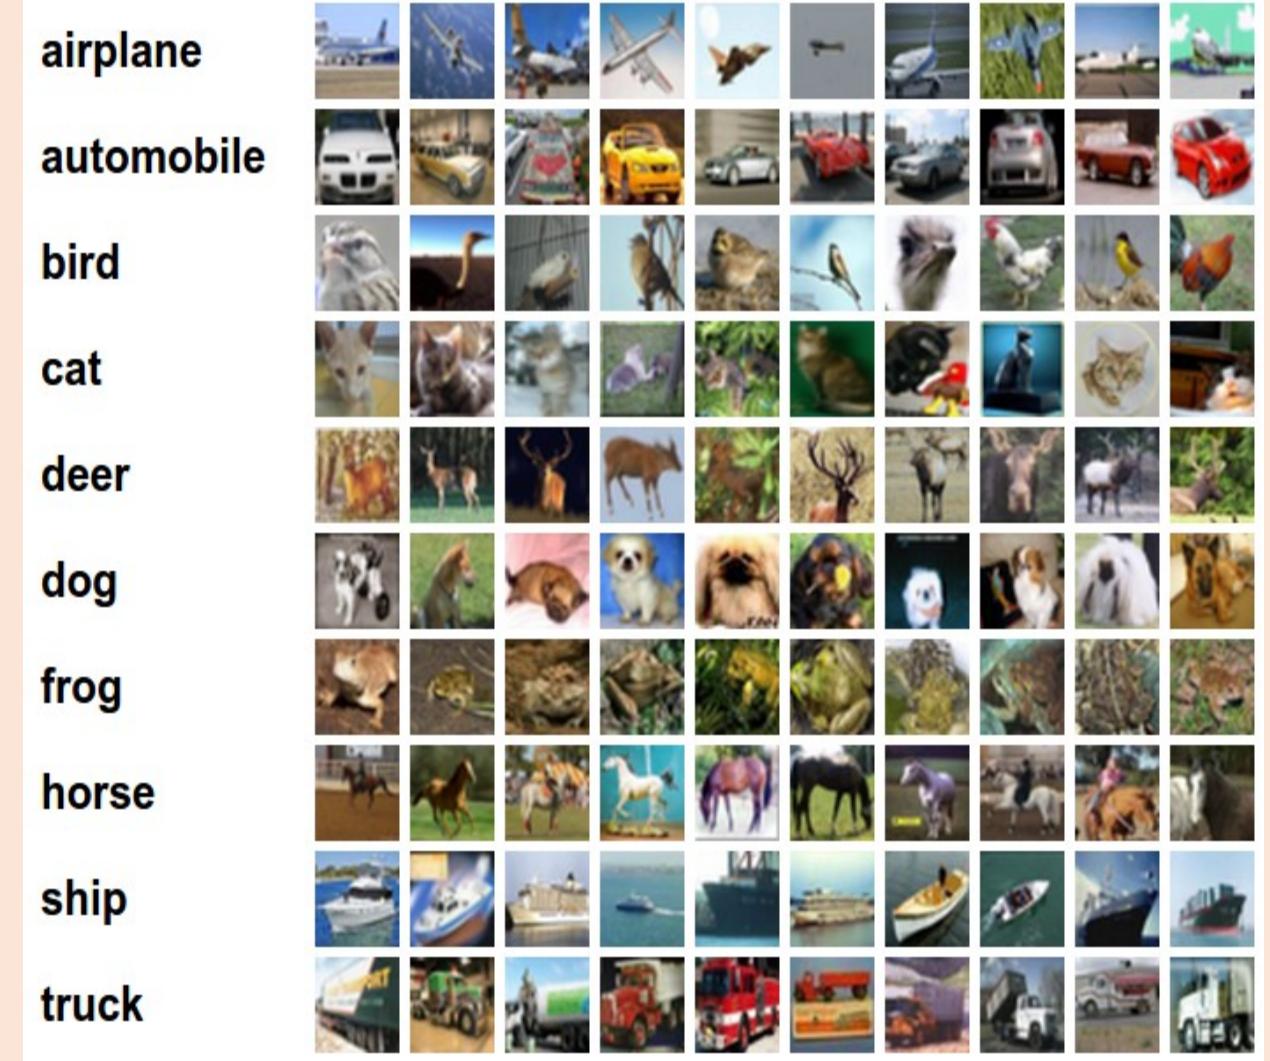
[0-1 - loss] 
$$L(\hat{y}) = \sum_i \chi_{\hat{y}_i \neq l}$$

[likelihood] 
$$L(\hat{y}) = m \cdot \log(p) + (k - m) \cdot \log(1 - p)$$

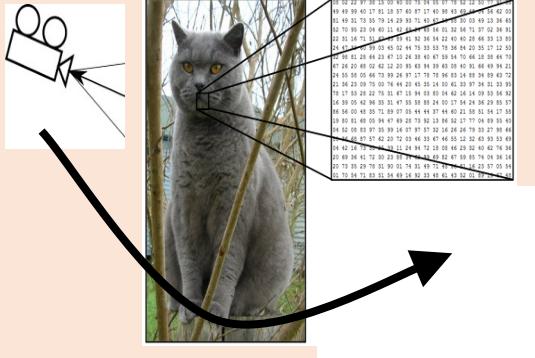
# Image Classification

## CIFAR-10

- 10 labels
- 50,000 training images
- each image is 32x32x3
- 10,000 test images.



# Image Classification

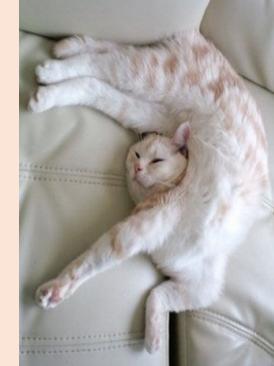


Camera pose

Illumination



Deformation



Occlusion



Background clutter



Intraclass variation



# Image Classification



image      parameters

$$f(\mathbf{x}, \mathbf{W})$$

classification

**[32x32x3]**  
array of numbers 0...1

# Image Classification

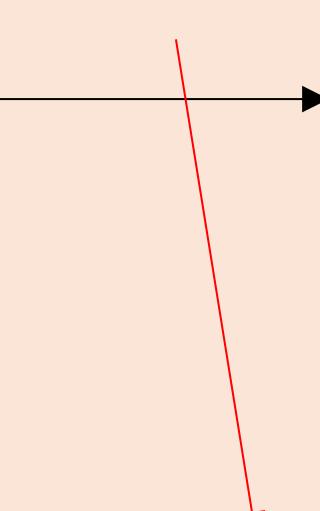


[32x32x3]  
array of numbers 0...1

$$f(x, W) = \boxed{W} \boxed{x}$$

10x1                    10x3072

3072x10



10 numbers,  
indicating class  
scores

parameters, or “weights”

# Image Classification



[32x32x3]  
array of numbers 0...1

$$f(x, W) = \boxed{W} \boxed{x}$$

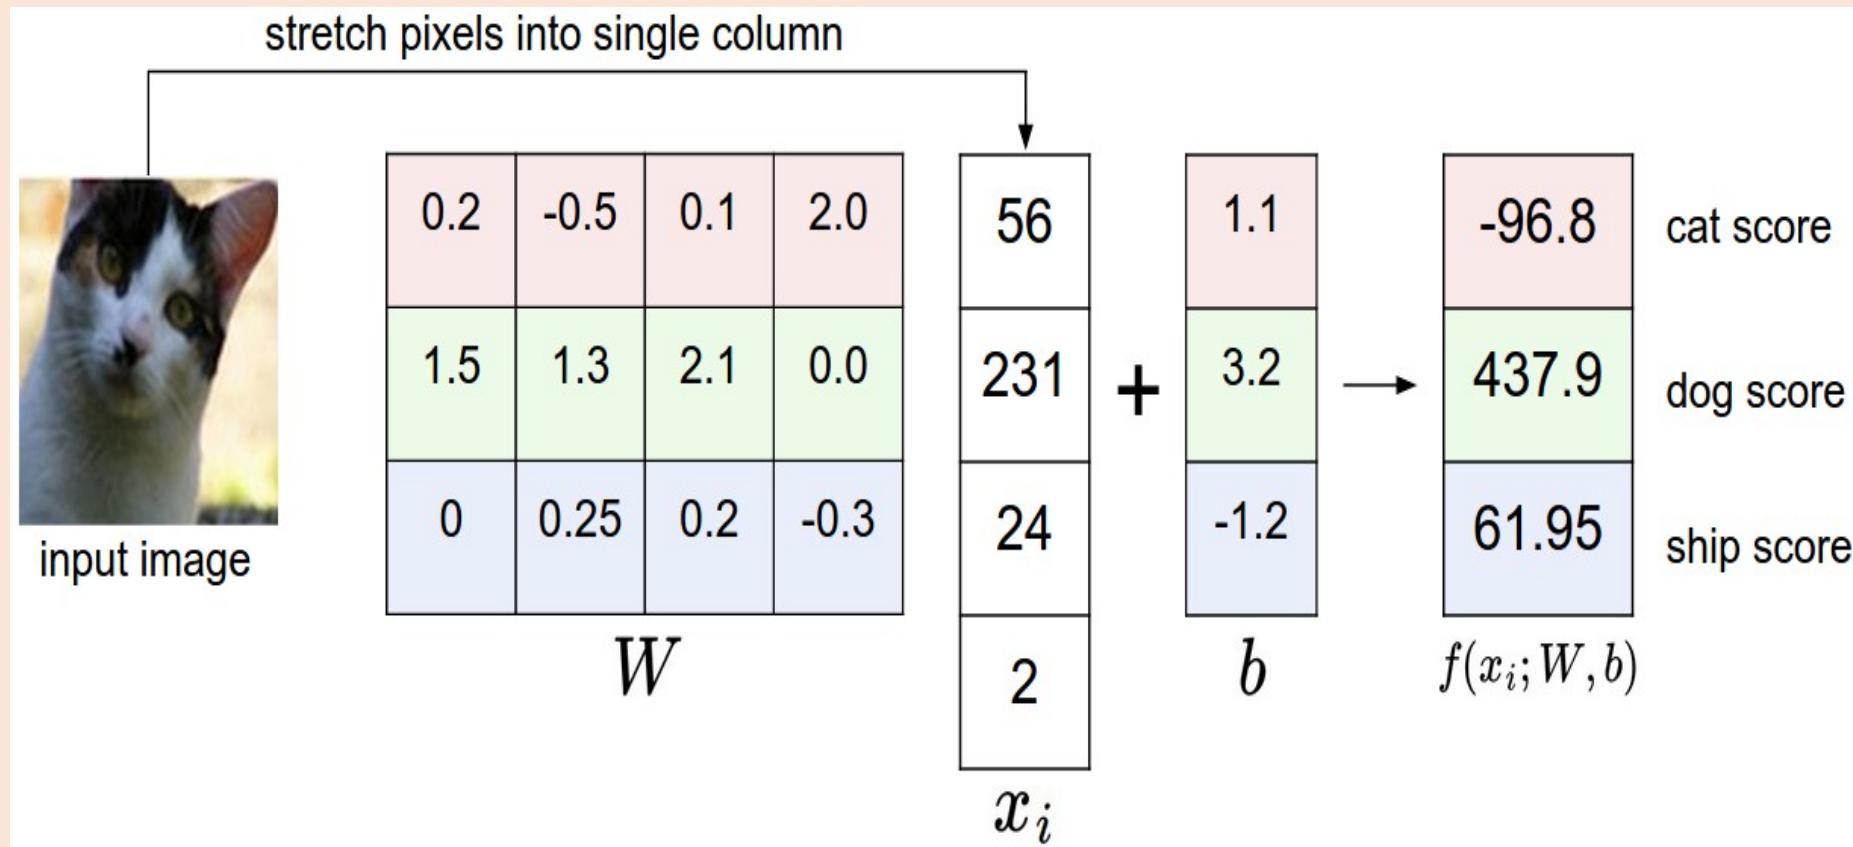
10x1                    10x 3072

3072x1                    (+b) 10x1

10 numbers,  
indicating class  
scores

parameters, or “weights”

# Geometric Interpretation



# Image Classification

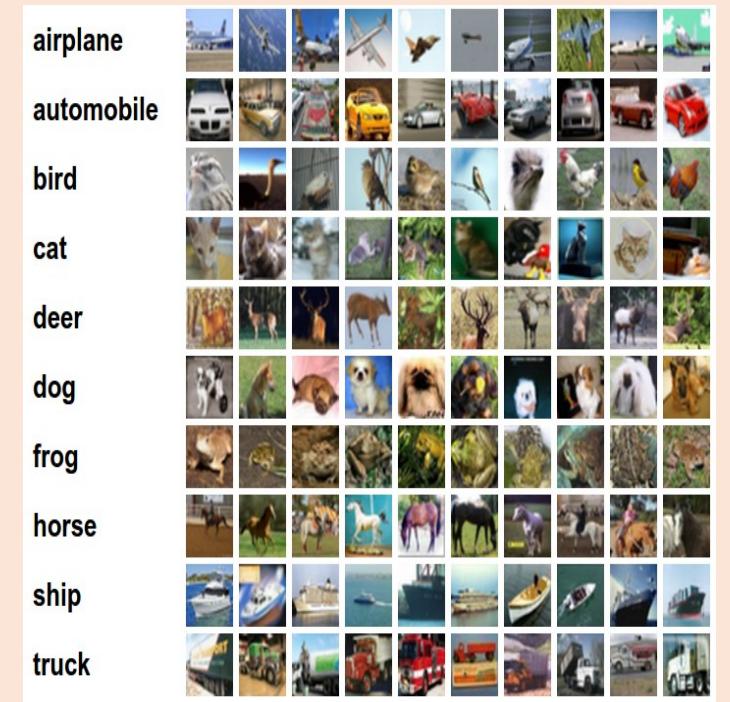
$$f(x_i, W, b) = Wx_i + b$$

Q: what does the linear classifier do?

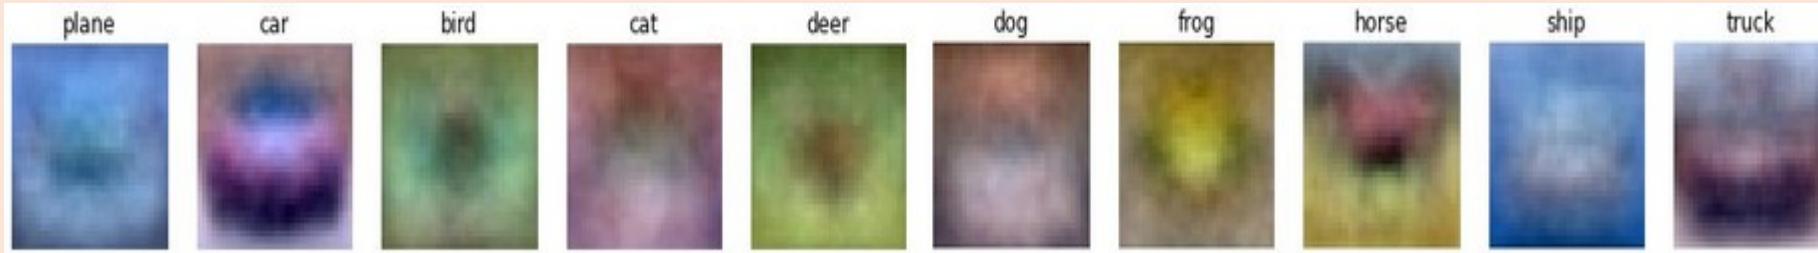


# Image Classification

Example trained weights of a linear classifier trained on CIFAR-10:



$$f(x_i, W, b) = Wx_i + b$$

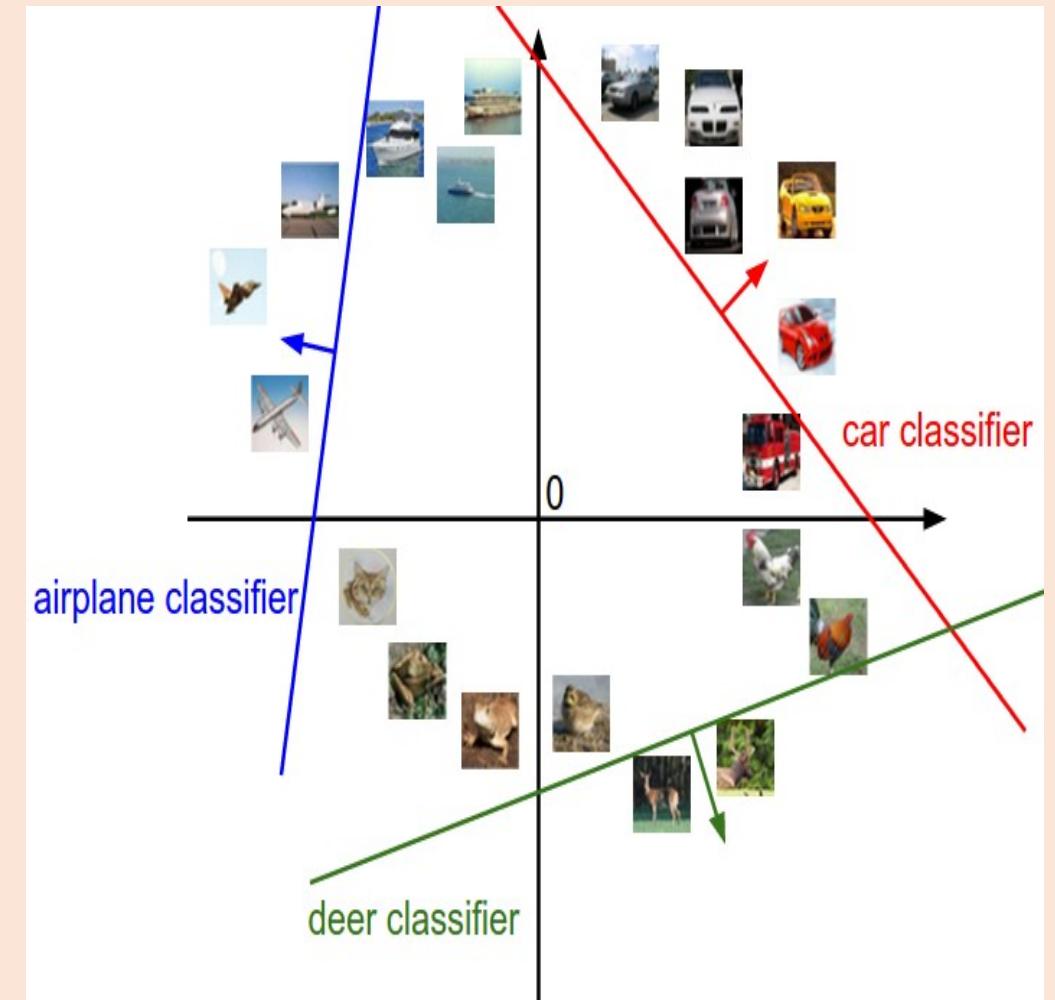


# Image Classification

[32x32x3]  
array of numbers 0...1  
(3072 numbers total)



$$f(x_i, W, b) = Wx_i + b$$



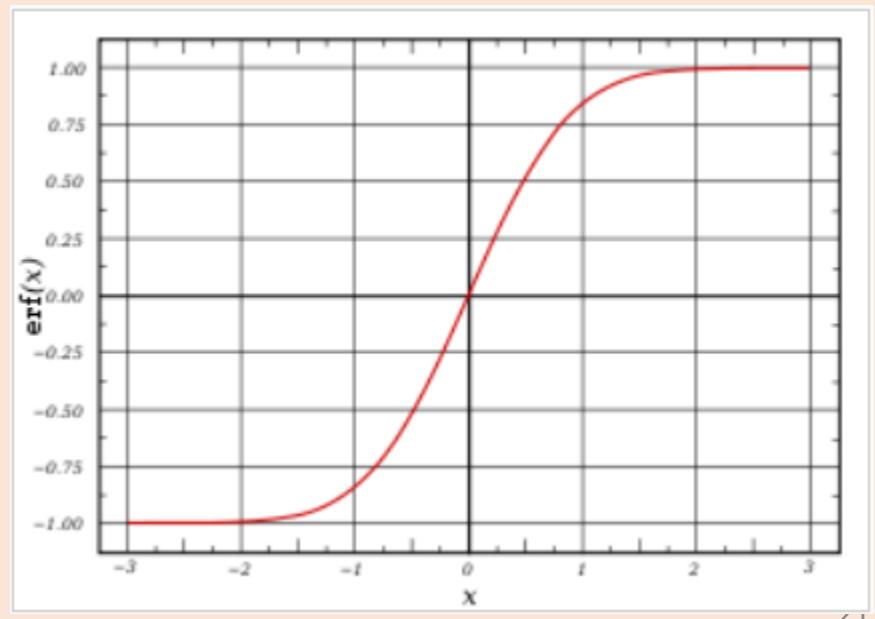
# Logistic Regression

$$P(Y = 1 | \mathbf{X}) = \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}}$$

Let  $\mathbf{X}$  be the data instance, and  $Y$  be the class label: Learn  $P(Y|\mathbf{X})$  directly

Let  $\mathbf{W} = (W_1, W_2, \dots, W_n)$ ,  $\mathbf{X} = (X_1, X_2, \dots, X_n)$ ,  $\mathbf{w}\mathbf{x}$  is the dot product

Sigmoid function:



# Logistic Regression

In logistic regression, we learn the conditional distribution  $P(y|x)$

Let  $p_y(x;w)$  be our estimate of  $P(y|x)$ , where  $w$  is a vector of adjustable parameters.

Assume there are two classes,  $y = 0$  and  $y = 1$  and

$$p_0(\mathbf{x}; \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}}$$

$$p_1(\mathbf{x}; \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}}$$

This is equivalent to

$$\log \frac{p_1(\mathbf{x}; \mathbf{w})}{p_0(\mathbf{x}; \mathbf{w})} = \mathbf{w}\mathbf{x}$$

That is, the log odds of class 1 is a linear function of  $\mathbf{x}$

# Learning Algorithm

The conditional data likelihood is the probability of the observed Y values in the training data, conditioned on their corresponding X values. We choose parameters w that satisfy

$$\mathbf{w} = \arg \max_{\mathbf{w}} \prod_l P(y^l | \mathbf{x}^l, \mathbf{w})$$

where  $\mathbf{w} = \langle w_0, w_1, \dots, w_n \rangle$  is the vector of parameters to be estimated,  $y^l$  denotes the observed value of Y in the l th training example, and  $\mathbf{x}^l$  denotes the observed value of X in the l th training example

# Implementation

```
In [9]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import h5py  
import scipy  
from scipy import ndimage  
from lr_utils import load_dataset
```

# Implementation

```
In [2]: #Load the data (cat/not cat datasets)
train_set_x_orig, train_set_y, test_set_x_orig, test_set_y, classes = load_dataset()

#Lets visualize the train set
print(train_set_x_orig)
print(train_set_y)
print(classes)
```

```
[[[ 17  31  56]
 [ 22  33  59]
 [ 25  35  62]
 ...
 [  1  28  57]
 [  1  26  56]
 [  1  22  51]]

 [[ 25  36  62]
 [ 28  38  64]
 [ 30  40  67]
 ...
 [  1  27  56]
 [  1  25  55]
 [  2  21  51]]

 [[ 32  40  67]
 [ 34  42  69]
 [ 35  42  70]]]
```

# Implementation

```
In [3]: #Lets get some basic data about our image numpy arrays  
m_train = train_set_x_orig.shape[0]  
m_test = test_set_x_orig.shape[0]  
num_px = train_set_x_orig.shape[1]  
  
print("Number of training examples: m_train = " + str(m_train))  
print("Number of test examples: m_test = " + str(m_test))  
print("Height/Width of each image: num_px = " + str(num_px))  
print("Each image is of size: (" + str(num_px) + ", " + str(num_px) + ", 3)")  
print("train_set_x shape: " + str(train_set_x_orig.shape))  
print("train_set_y shape: " + str(train_set_y.shape))  
print("test_set_x shape : " + str(test_set_x_orig.shape))  
print("test_set_y shape: " + str(test_set_y.shape))
```

```
Number of training examples: m_train = 209  
Number of test examples: m_test = 50  
Height/Width of each image: num_px = 64  
Each image is of size: (64, 64, 3)  
train_set_x shape: (209, 64, 64, 3)  
train_set_y shape: (1, 209)  
test_set_x shape : (50, 64, 64, 3)  
test_set_y shape: (1, 50)
```

# Implementation

```
In [4]: #Will now flatten the numpy array from (num_px, num_px, 3) to (num_px*num_px*3, 1)
#this will make it easier for us so that each image in one numpy array column
train_set_x_flatten = train_set_x_orig.reshape(train_set_x_orig.shape[0], -1).T
test_set_x_flatten = test_set_x_orig.reshape(test_set_x_orig.shape[0], -1).T

print("train_set_x_flatten shape: " + str(train_set_x_flatten.shape))
print("train_set_y shape: " + str(train_set_y.shape))
print("test_set_x_flatten shape: " + str(test_set_x_flatten.shape))
print("test_set_y shape: " + str(test_set_y.shape))
```

```
#Standardize the dataset for images by dividing each by 255
```

```
train_set_x = train_set_x_flatten/255
test_set_x = test_set_x_flatten/255
```

```
train_set_x_flatten shape: (12288, 209)
train_set_y shape: (1, 209)
test_set_x_flatten shape: (12288, 50)
test_set_y shape: (1, 50)
```

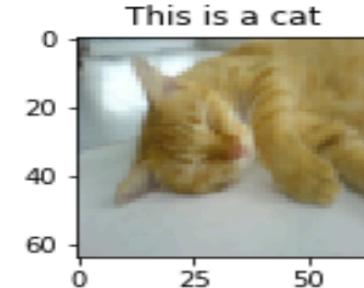
# Implementation

```
In [5]: ROWS=64;COLS=64;CHANNELS=3;
classes = {0:'not a cat',
           1:'a cat'}

def show_image(X, y, idx):
    image = X[idx]
    image = image.reshape((ROWS, COLS, CHANNELS))
    plt.figure(figsize=(4,2))
    plt.imshow(image)
    plt.title("This is {}".format(classes[y[0],idx]))
    plt.show()

def show_image_prediction(X, idx, model):
    image = X[idx].reshape(1, -1)
    image_class = classes[model.predict(image).item()]
    image = image.reshape((ROWS, COLS, CHANNELS))
    plt.figure(figsize=(4,2))
    plt.imshow(image)
    plt.title("Test {}: I think this {}".format(idx, image_class))
    plt.show()
```

```
In [13]: x=train_set_x_orig
y=train_set_y
show_image(x, y, 121)
print(x[121].shape)
print(x[121])
```



```
(64, 64, 3)
[[[106 102 96]
  [90 97 103]
  [104 122 124]
  ...
  [169 148 104]
  [169 147 97]
  [168 145 92]]
 [[117 110 99]
  [108 114 112]
  [119 126 125]
  ...
  [171 149 104]
  [171 149 103]]]
```

# Implementation

```
In [14]: #We will be using a sigmoid function for our Activation, in Neural Networks most are not ReLU due to computational cost  
  
def sigmoid(z):  
    s = 1/(1+np.exp(-(z)))  
    return s  
  
#Create function to set both w and b to 0 to start with  
def initialize_with_zeros(dim):  
    w = np.zeros((dim,1))  
    b = 0  
    return w,b
```

# Implementation

```
def predict(w, b, X):
    m = X.shape[1]
    Y_prediction = np.zeros((1,m))
    w = w.reshape(X.shape[0],1)
    A = sigmoid(np.dot(w.T,X)+b)
    for i in range(A.shape[1]):
        if A[0,i] <= 0.5:
            Y_prediction[0,i] = 0
        else:
            Y_prediction[0,i] = 1
    return Y_prediction
```