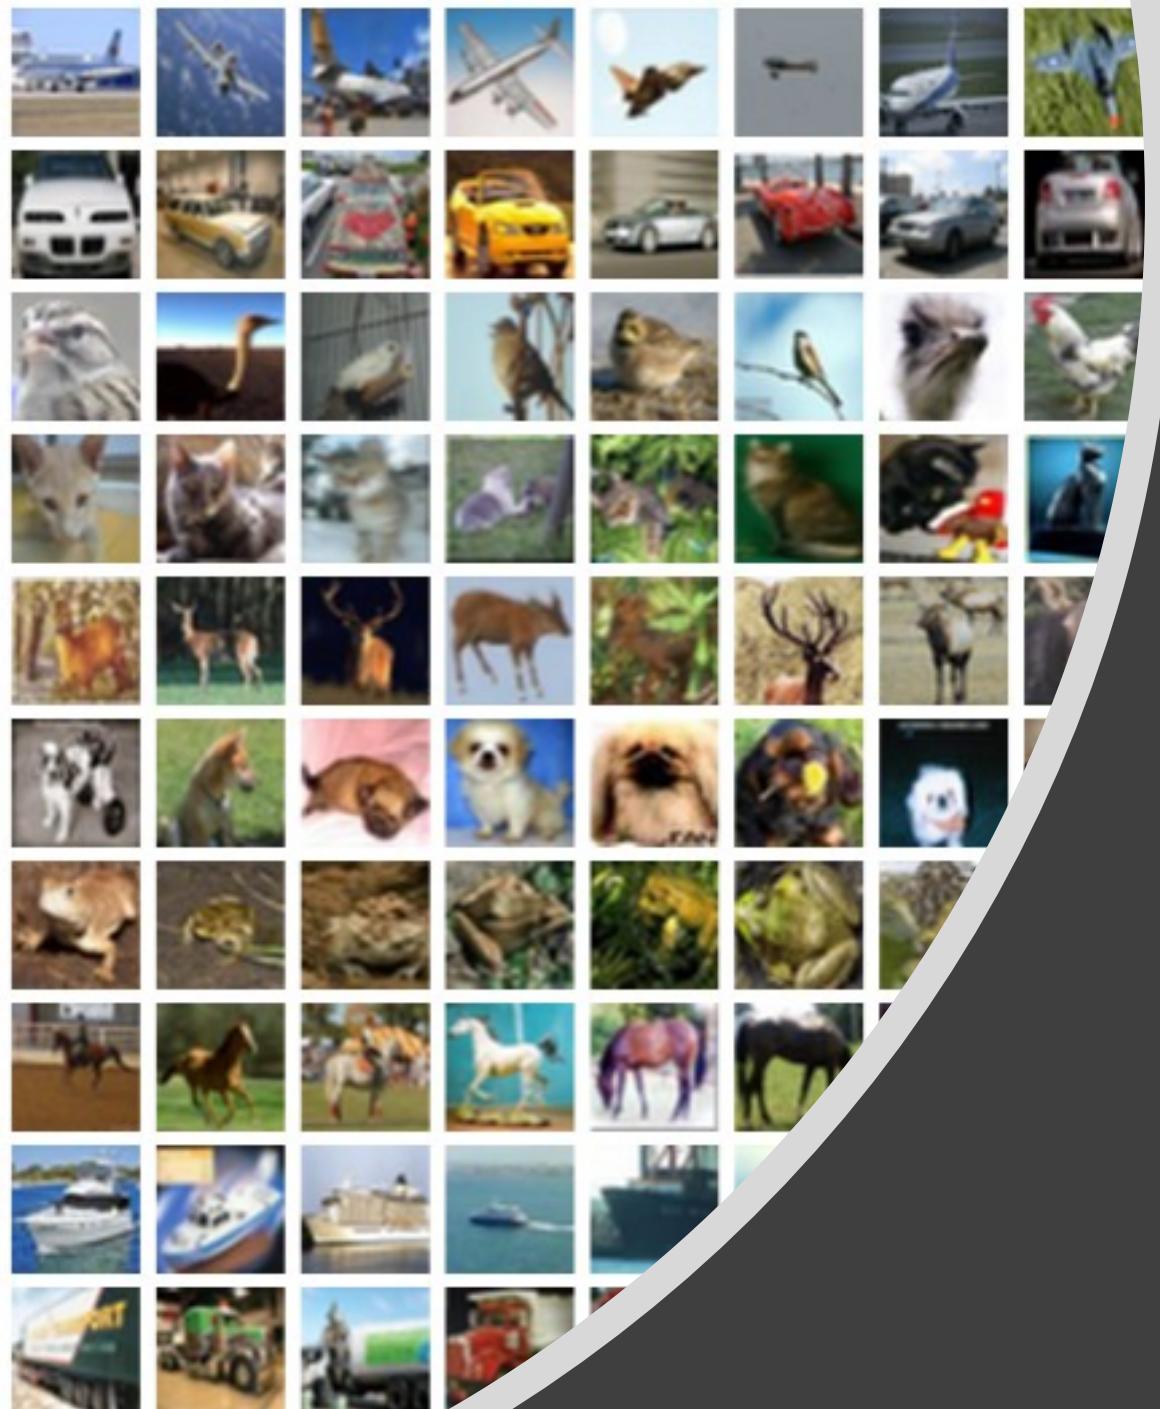


le



Dynamic Learning

Perceptron

Syllabus Update

Week 10	logistic regression, neural networks and gradient descent	Initial draft of project proposal due
Week 11	Overfitting and avoidance	Exercise 3 due
Week 12	Model Evaluation	Final draft of project proposal due
Week 13	Bayesian Methods	Exercise 4 due
Week 14	Course Recap	
Week 15	FINAL PROJECTS AND PRESENTATIONS DUE	

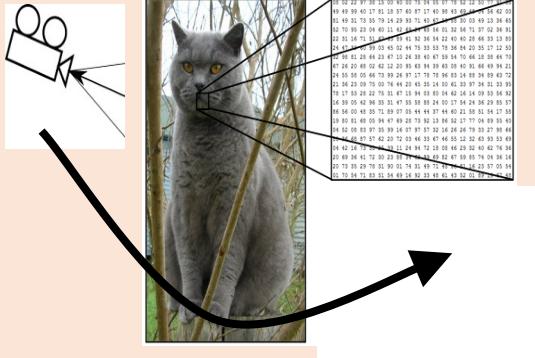
Final Project

- Open ended any topic involving learning models we will encounter by end of course
 - Decision trees
 - Logistic regression
 - Neural networks
- Responsibility:
 - Identify a problem you would like to work on
 - Identify source of data
 - Write the scripts to ingest the data
 - identify an appropriate Machine learning library
 - Write the scripts for passing the data to the libraries
 - Interpret the results

Lecture Plan

- Linear Classifiers
- Perceptron
 - Perceptron learning
- Gradient Descent
 - Stochastic Gradient Descent

Image Classification

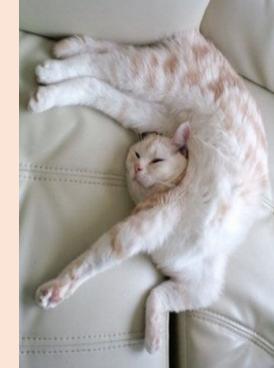


Camera pose

Illumination



Deformation



Occlusion



Background clutter



Intraclass variation



Image Classification



image parameters

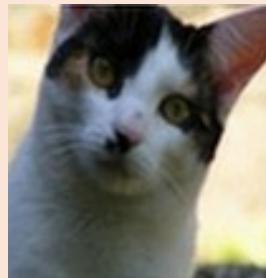
$$f(\mathbf{x}, \mathbf{W})$$

classification

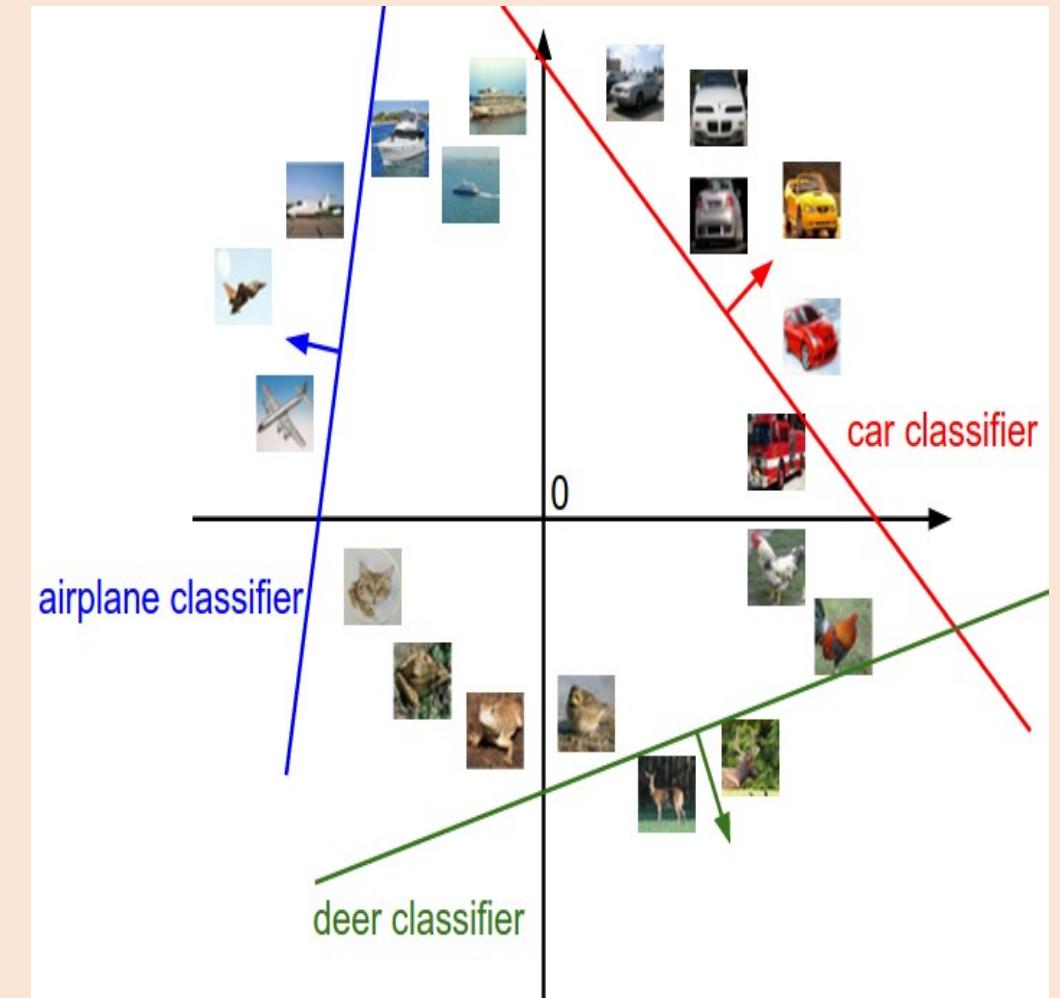
$$f(x_i, W, b) = Wx_i + b$$

Image Classification

[32x32x3]
array of numbers 0...1
(3072 numbers total)



$$f(x_i, W, b) = Wx_i + b$$

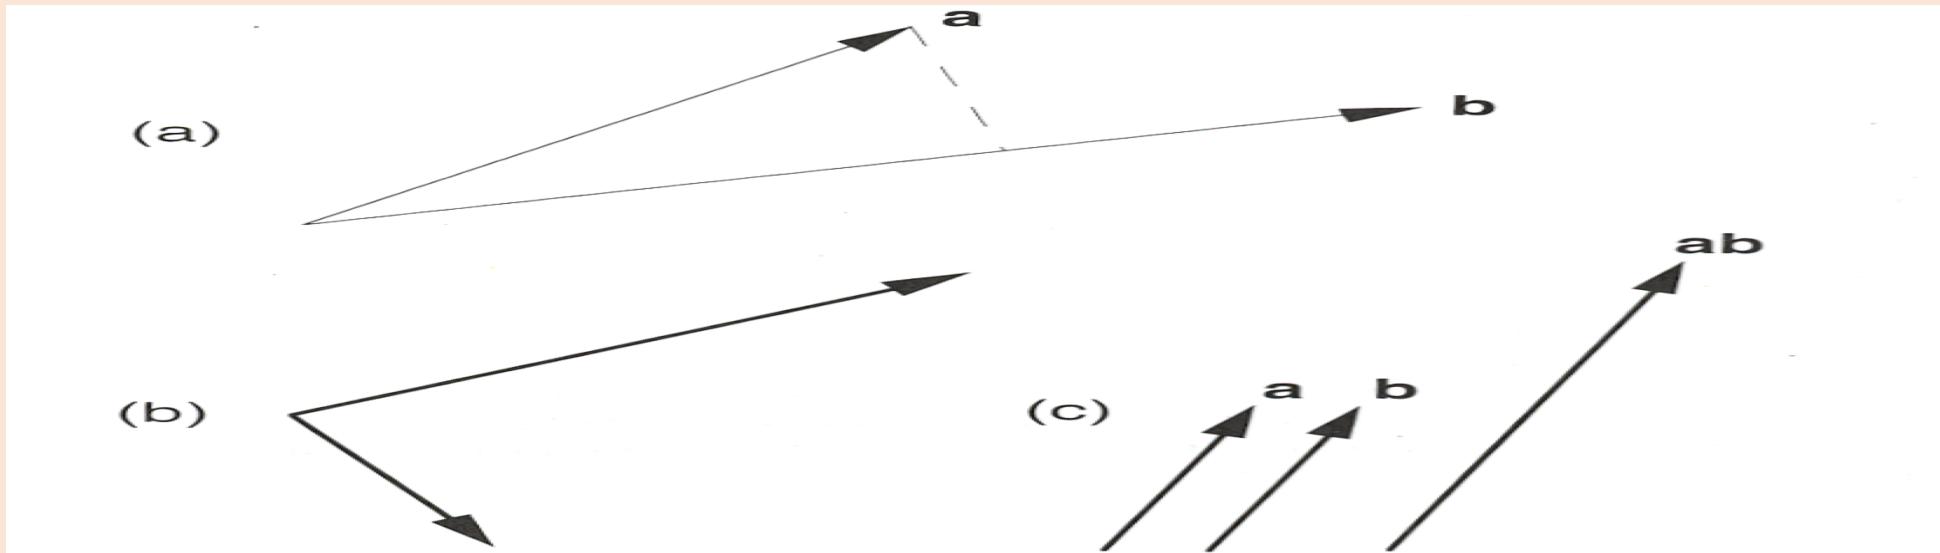


Inner Product

$$\langle w, x \rangle = \sum_i w_i \cdot x_i$$

$$n = \langle w, x \rangle = ||w|| \cdot ||x|| \cdot \cos(\theta)$$

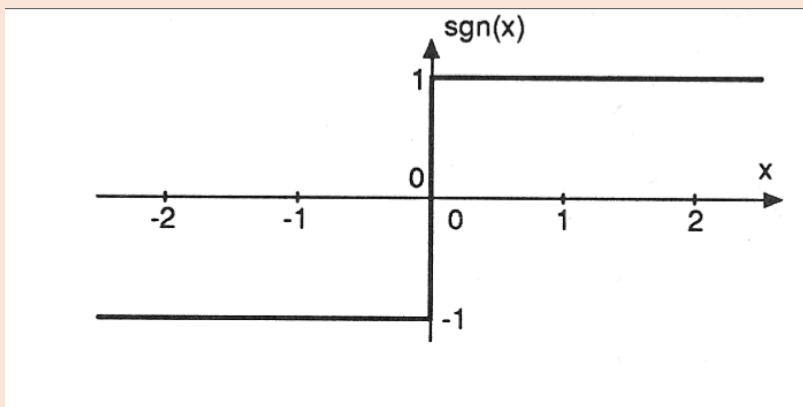
- A measure of the projection of one vector onto another



Activation Function

$$f(x) = \operatorname{sgn} \left(\sum_i w_i \cdot x_i \right)$$

$$\operatorname{sgn}(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ -1 & \text{if } x < 0 \end{cases}$$

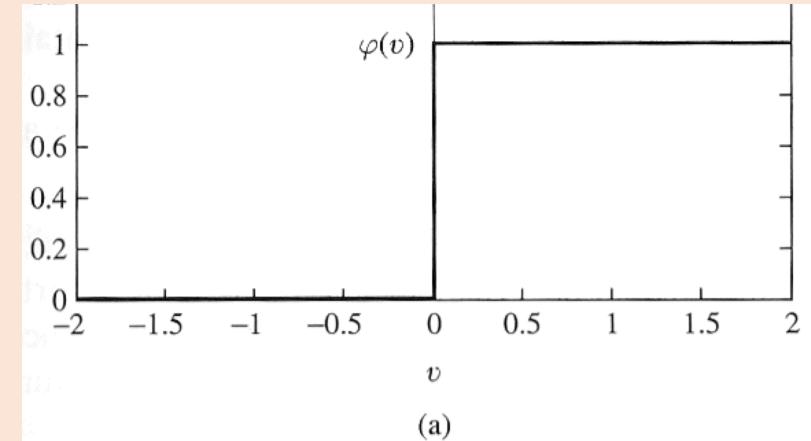


Activation Function

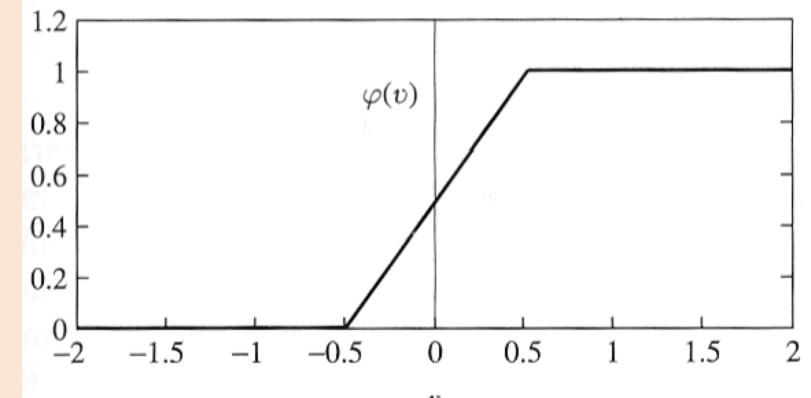
$$\varphi(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ 0 & \text{if } x < 0 \end{cases}$$

$$\varphi(x) = \begin{cases} 1 & \text{if } x > 1/2 \\ x & \text{if } -\frac{1}{2} < x \leq 1/2 \\ -1 & \text{if } x < -1/2 \end{cases}$$

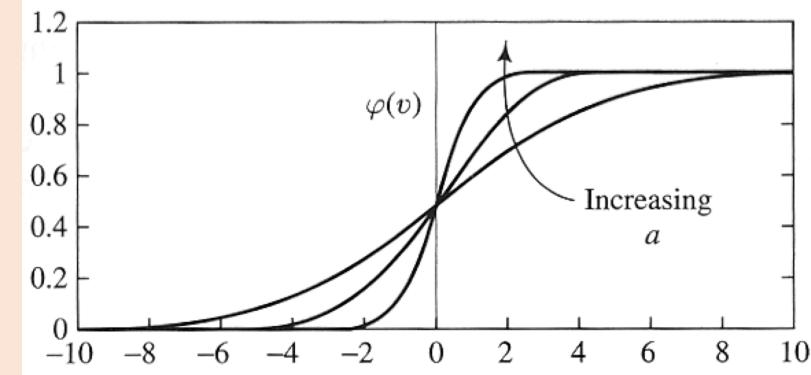
$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



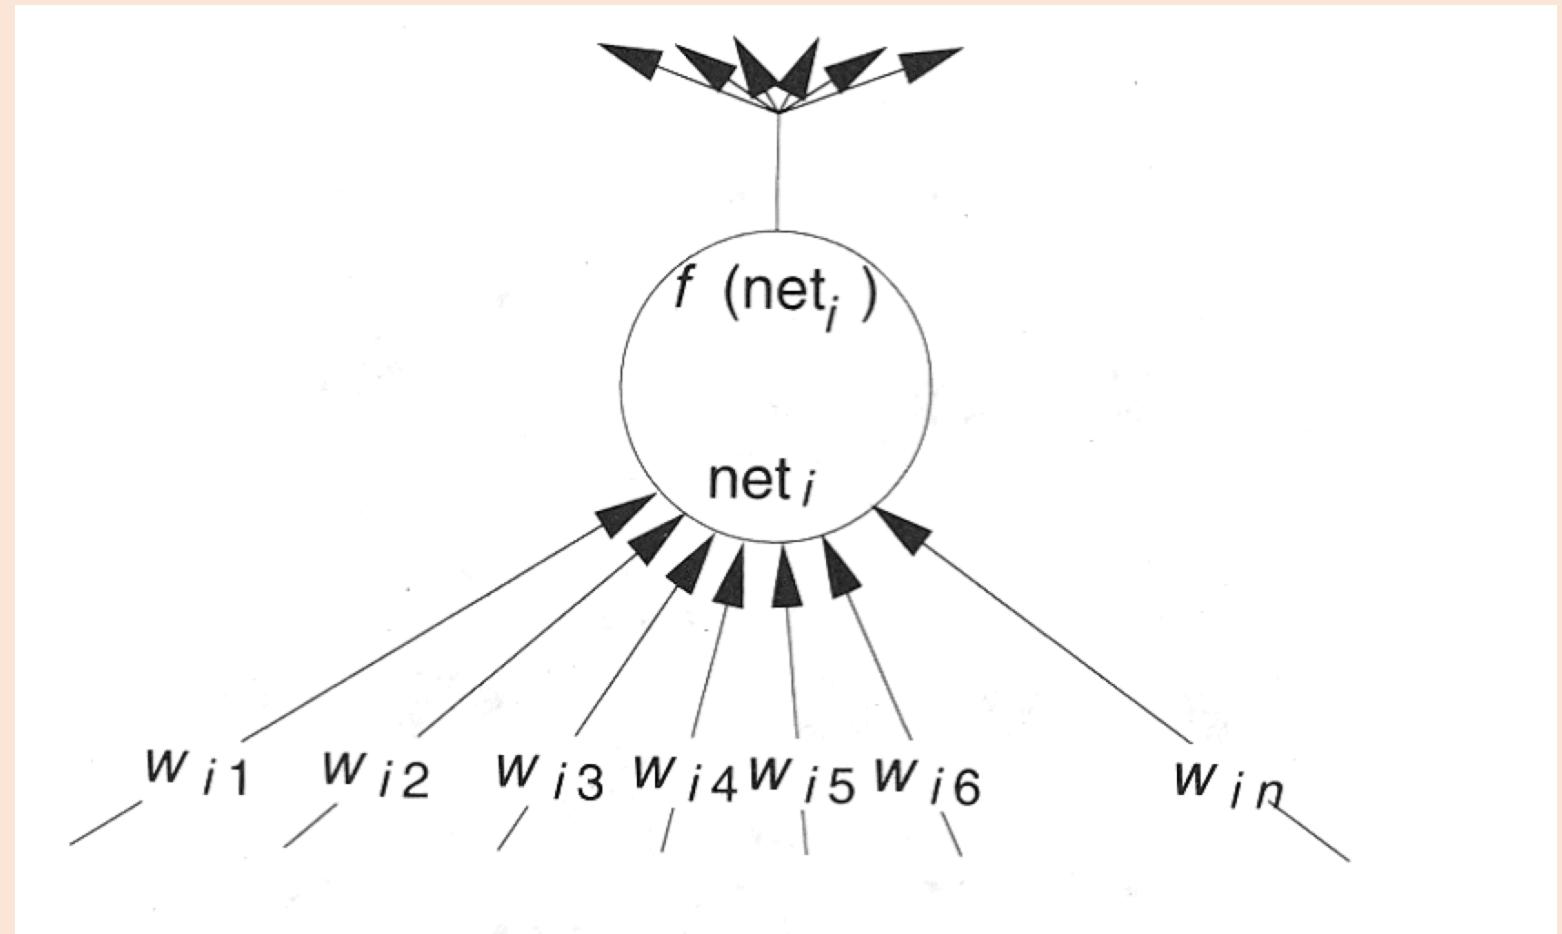
(a)



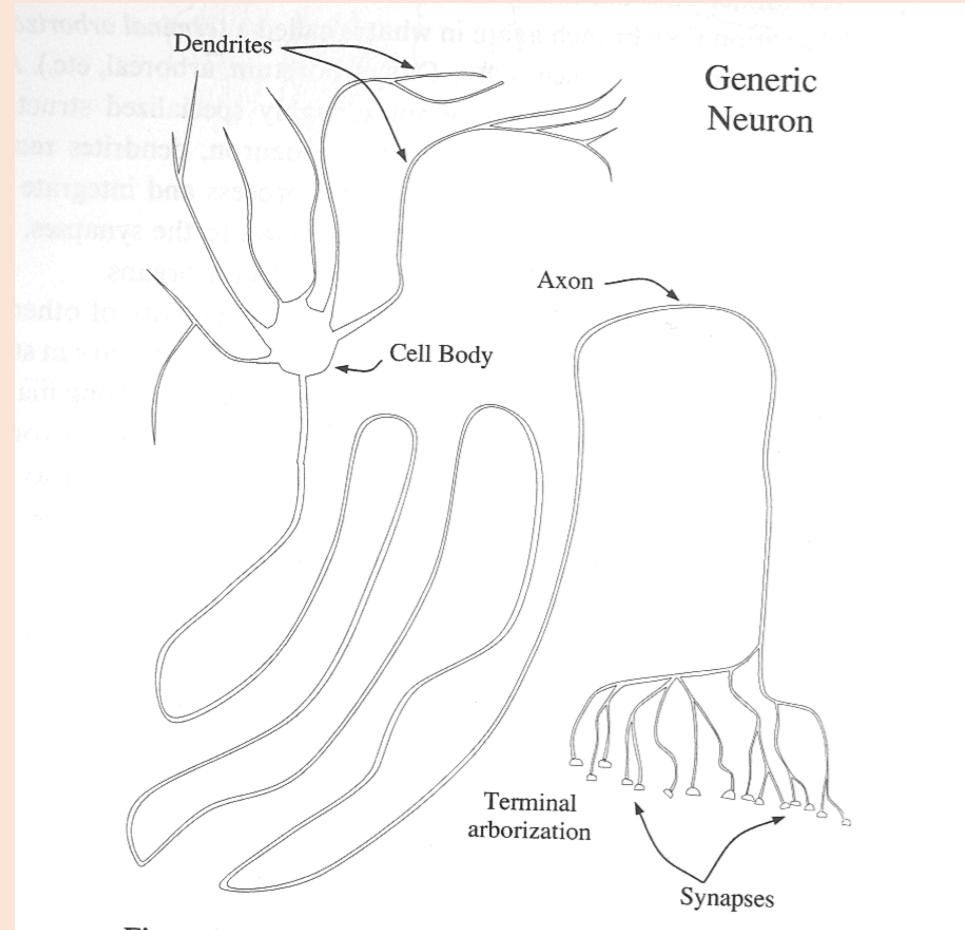
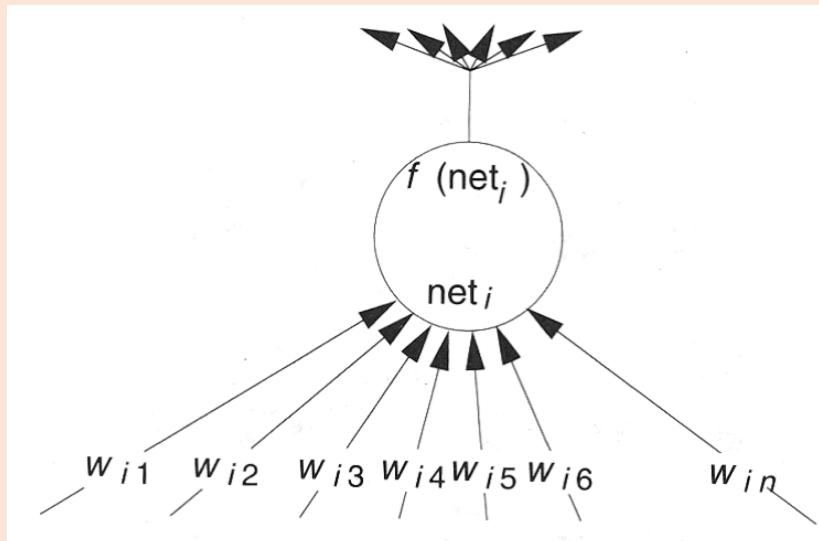
(b)



Graphical Representation

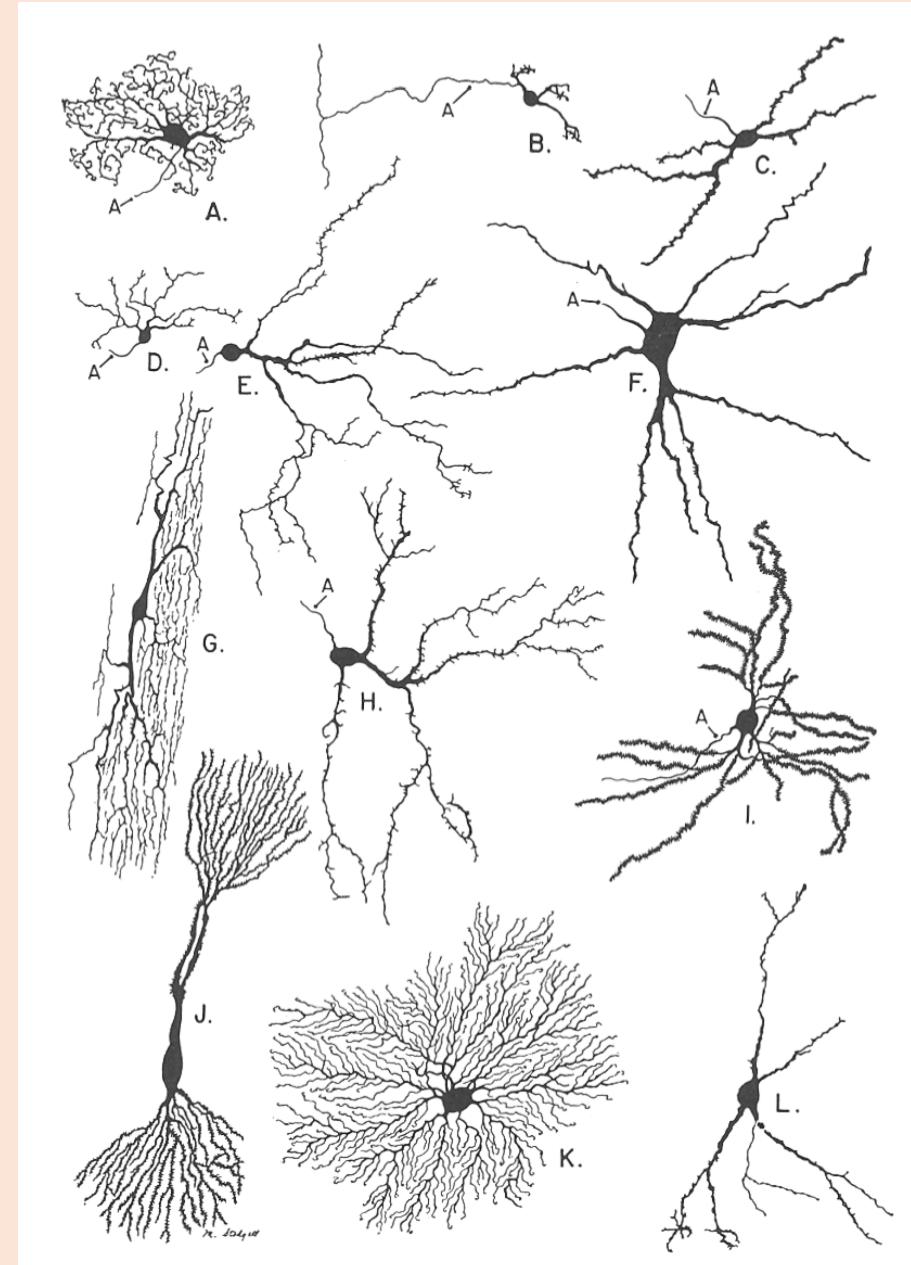


Resemblance to real neurons

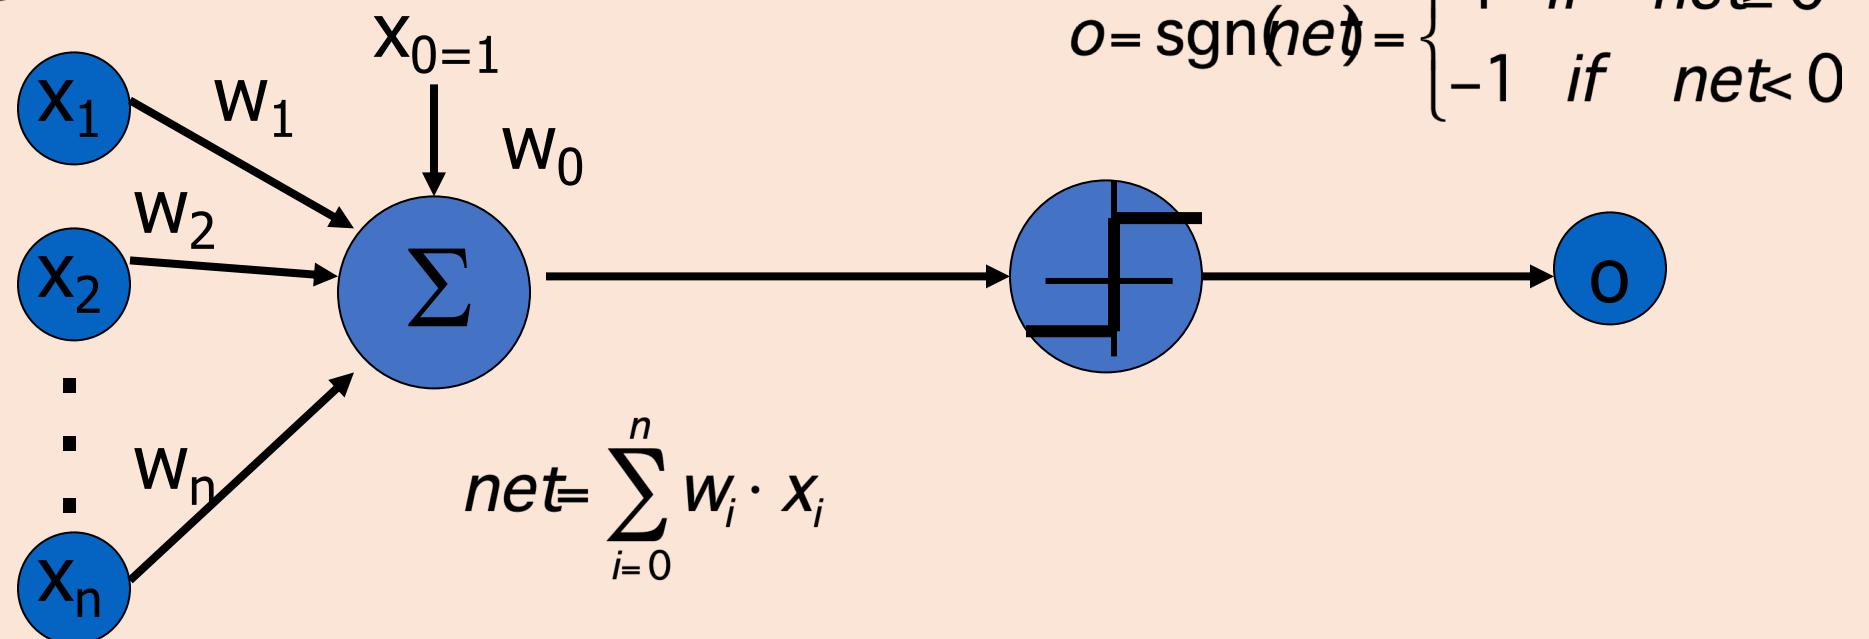


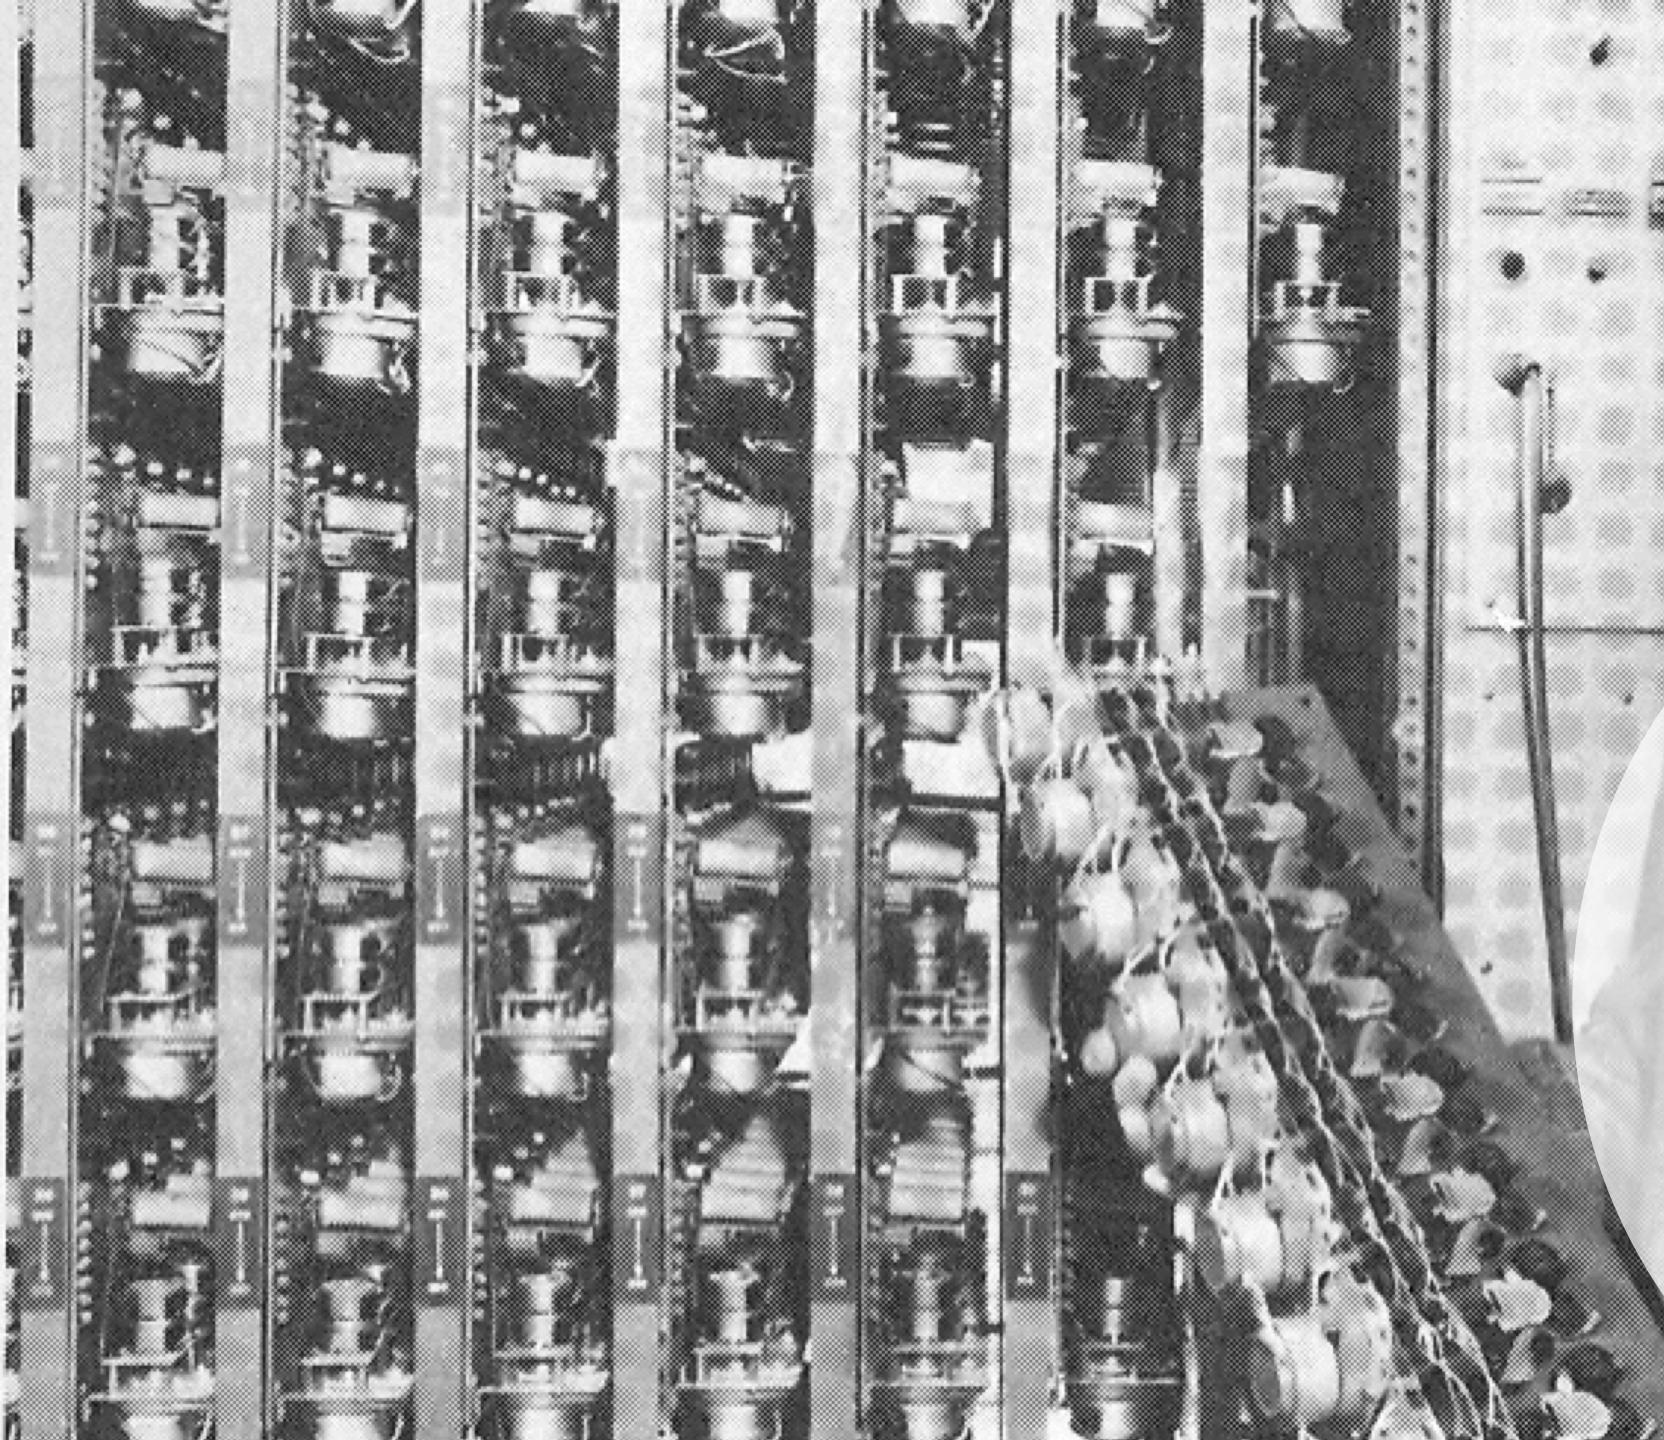
Resemblance to real neurons

- 10^{40} Neurons
- 10^{4-5} connections per neuron



Perceptron





Frank Rosenblatt

1928 - 1969

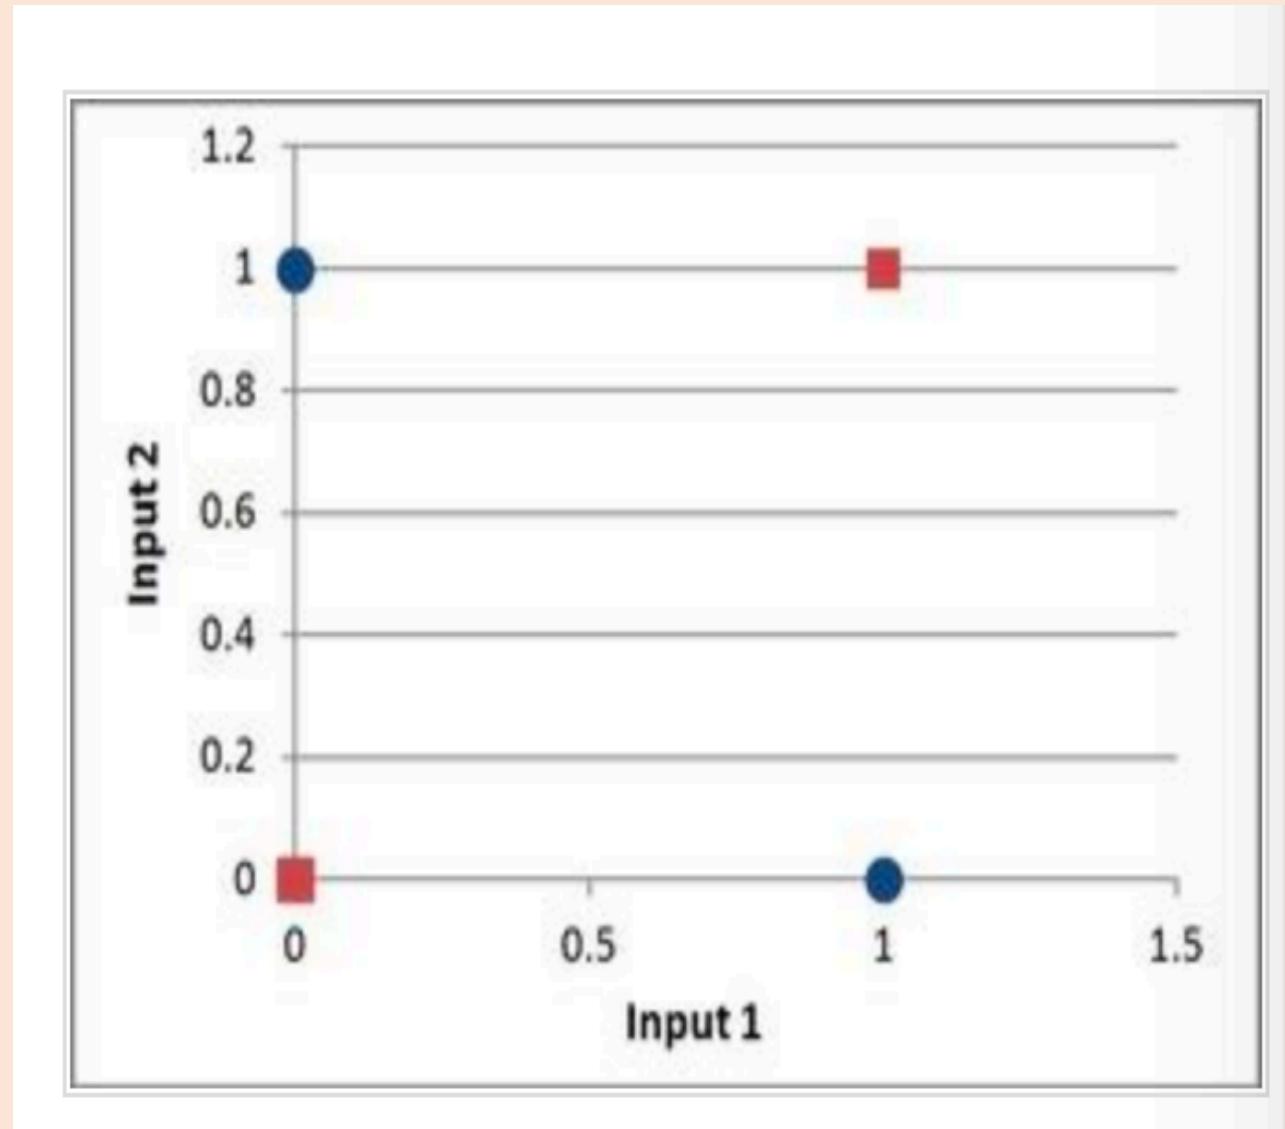
Frank Rosenblatt

- Rosenblatt's bitter rival and professional nemesis was Marvin Minsky of Carnegie Mellon University
- Minsky despised Rosenblatt, hated the concept of the perceptron, and wrote several polemics against him
- For years Minsky crusaded against Rosenblatt on a very nasty and personal level, including contacting every group who funded Rosenblatt's research to denounce him as a charlatan, hoping to ruin Rosenblatt professionally and to cut off all funding for his research in neural nets

XOR Problem

Minsky and Papert in
mid 1960

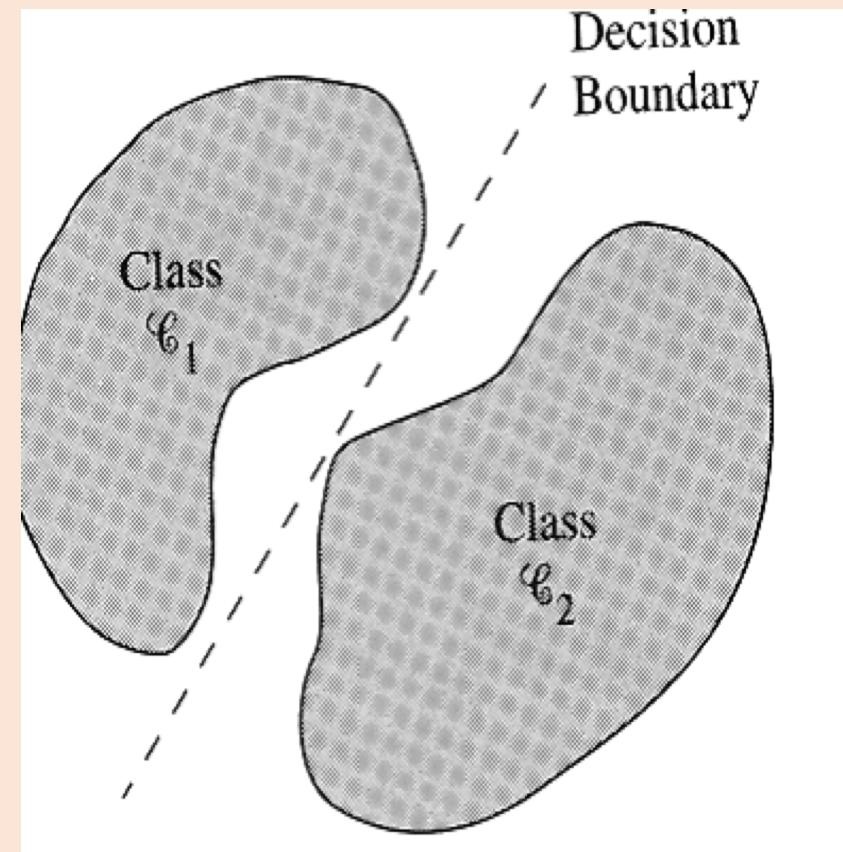
	input 1	input 2	desired output
object 1	0	0	0
object 2	0	1	1
object 3	1	0	1
object 4	1	1	0



Resemblance to real neurons

The goal of a perceptron is to correctly classify the set of pattern $D=\{x_1, x_2, \dots, x_m\}$ into one of the classes C_1 and C_2

The output for class C_1 is $out=1$ and for C_2 is $out=-1$

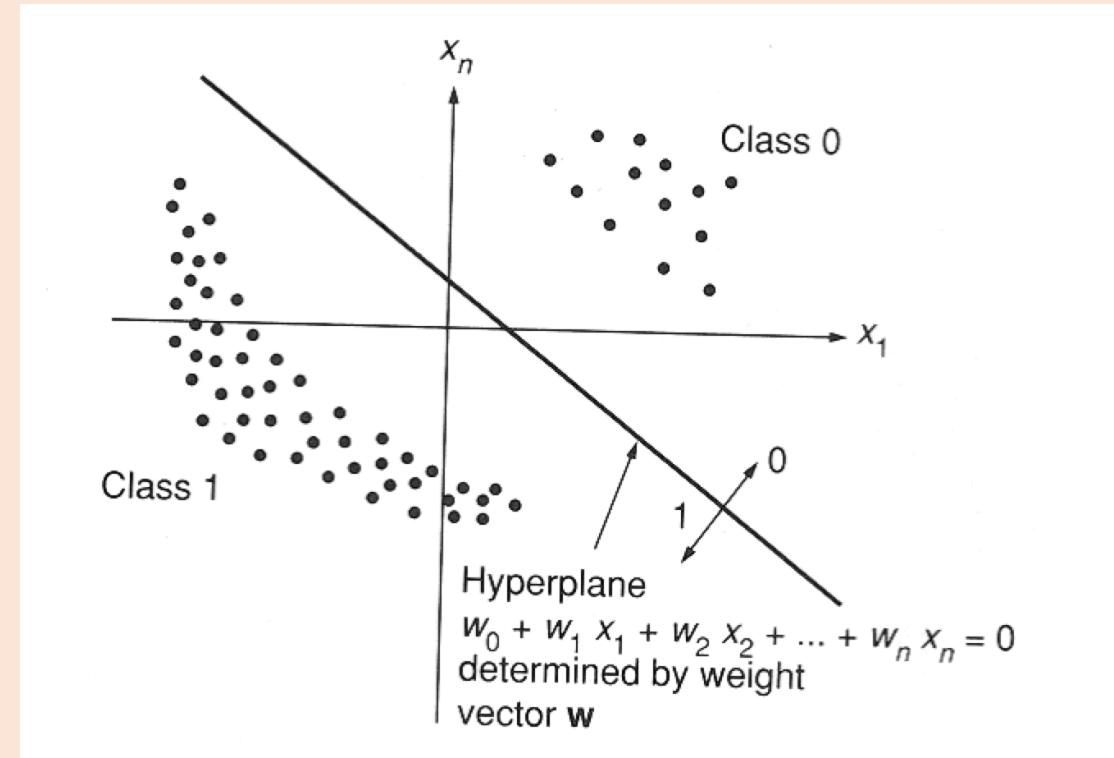


Linearly separable patterns

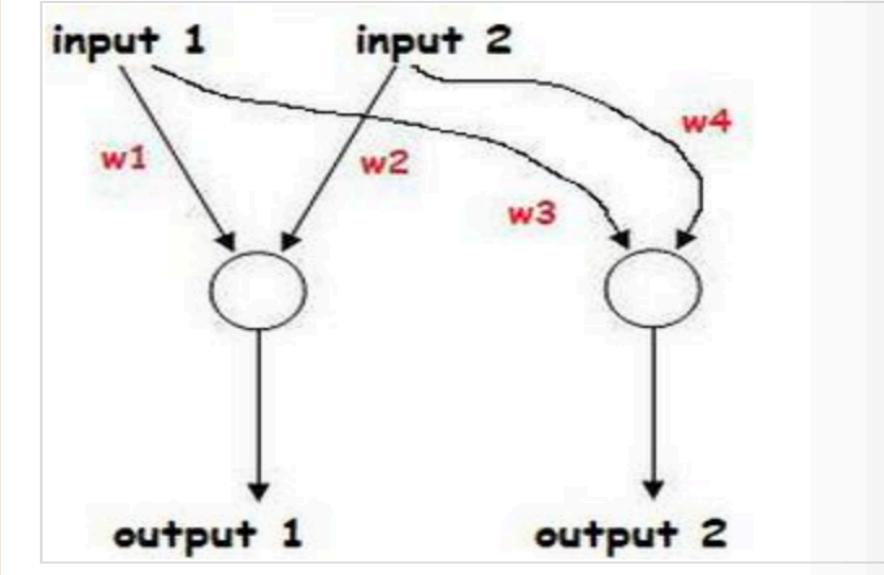
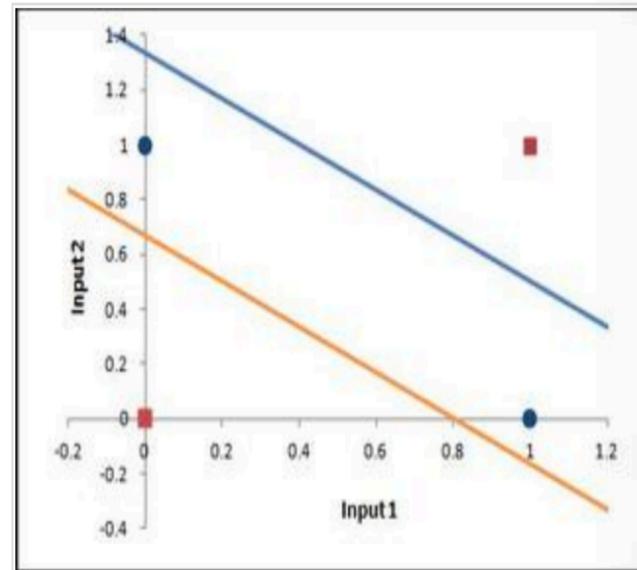
$$\sum_i w_i \cdot x_i \leq 0 \text{ for class 1}$$

$$out = sgn \left(\sum_i w_i \cdot x_i \right)$$

$$\sum_i w_i \cdot x_i > 0 \text{ for Class 0}$$



Multi level networks



	input 1	input 2	output
object 1	0	0	0
object 2	0	1	1
object 3	1	0	1
object 4	1	1	1

	input 1	input 2	output
object 1	0	0	0
object 2	0	1	0
object 3	1	0	0
object 4	1	1	1

	input 1	input 2	output
object 1	0	0	0
object 2	1	0	1
object 3	1	0	1
object 4	1	1	0

So we use the logic table that user just made for the

Perceptron learning rule

- Consider linearly separable problems
- How to find appropriate weights
- Look if the output pattern o belongs to the desired class, has the desired value d

$$W^{new} = W^{old} + \Delta W \quad \Delta W = \eta(d - o)$$

- η is called the **learning rate**
- $0 < \eta \leq 1$

Perceptron learning rule

- The algorithm converges to the correct classification
 - if the training data is linearly separable
 - and η is sufficiently small
- When assigning a value to η we must keep in mind two conflicting requirements
 - Averaging of past inputs to provide stable weights estimates, which requires small η
 - Fast adaptation with respect to real changes in the underlying distribution of the process responsible for the generation of the input vector x , which requires large η

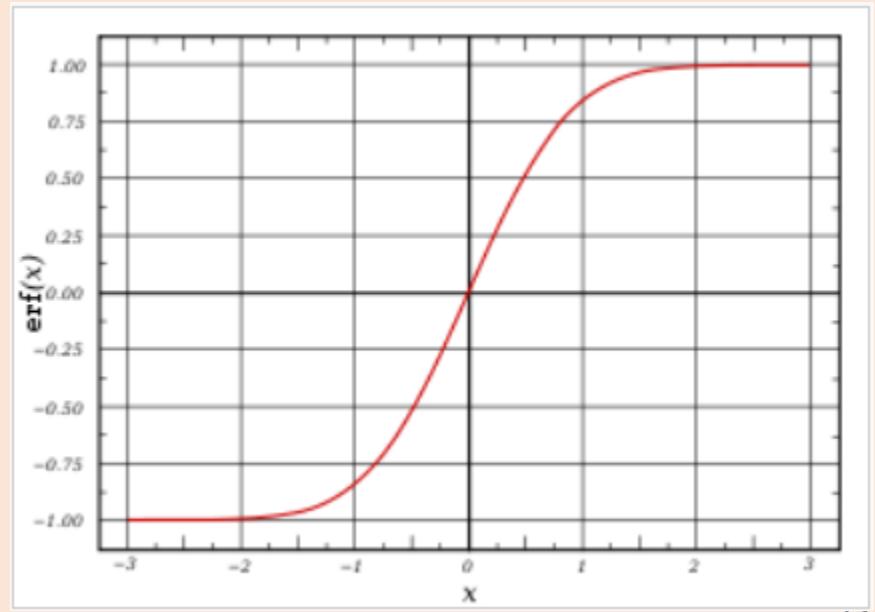
Logistic Regression

$$P(Y = 1|X) = \frac{1}{1 + e^{-\langle x, w \rangle}}$$

Let X be the data instance, and Y be the class label: Learn $P(Y|X)$ directly

Let $W = (W_1, W_2, \dots, W_n)$, $X = (X_1, X_2, \dots, X_n)$, WX is the dot product

Sigmoid function:



Logistic Regression

In logistic regression, we learn the conditional distribution $P(y|x)$

Let $p_y(x;w)$ be our estimate of $P(y|x)$, where w is a vector of adjustable parameters.

Assume there are two classes, $y = 0$ and $y = 1$ and

$$p_0(\mathbf{x}; \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}}$$

$$p_1(\mathbf{x}; \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}}$$

This is equivalent to

$$\log \frac{p_1(\mathbf{x}; \mathbf{w})}{p_0(\mathbf{x}; \mathbf{w})} = \mathbf{w}\mathbf{x}$$

That is, the log odds of class 1 is a linear function of \mathbf{x}

Learning Algorithm

The conditional data likelihood is the probability of the observed Y values in the training data, conditioned on their corresponding X values. We choose parameters w that satisfy

$$\mathbf{w} = \arg \max_{\mathbf{w}} \prod_l P(y^l | \mathbf{x}^l, \mathbf{w})$$

where $\mathbf{w} = \langle w_0, w_1, \dots, w_n \rangle$ is the vector of parameters to be estimated, y^l denotes the observed value of Y in the l th training example, and \mathbf{x}^l denotes the observed value of X in the l th training example