


le



Dynamic Learning

Logistic Regression,
Neural Networks and
Gradient Descent

Final Project Proposal

- Project overview (3 paragraphs)
- Objectives
 - What is the purpose of the project
 - How will the project be evaluated
- Data to be used
- Learning model
- Libraries


Perceptron Algorithm

- Perceptron Algorithm
 - Multi layered perceptron
- Regression
 - Linear
 - Logistic
- Gradient Descent
 - Batch Gradient Descent
 - Stochastic Gradient Descent
- Introduction to Neural Networks
 - Feedforward Neural Networks

The Perceptron

$$g(z) = \text{sgn}(z) = \begin{cases} -1 & z \leq 0 \\ +1 & z > 0 \end{cases}$$

$$h_w(x) = g(\sum w_i x_i) = g(W \cdot X)$$



Perceptron Algorithm

Given a training set $(X_0, y_0), \dots (X_k, y_k)$

$$W \leftarrow (w_{0,1}, \dots w_{0,n})$$

for (X, y) *in* $(X_0, y_0), \dots (X_k, y_k)$:

$$h \leftarrow g(\sum w_i x_i)$$

$$w_i \leftarrow w_i + \eta \cdot x_i \cdot (y - h)$$

Perceptron learning rule

- The algorithm converges to the correct classification
 - if the training data is linearly separable
 - and η is sufficiently small
- When assigning a value to η we must keep in mind two conflicting requirements
 - Averaging of past inputs to provide stable weights estimates, which requires small η
 - Fast adaptation with respect to real changes in the underlying distribution of the process responsible for the generation of the input vector \mathbf{x} , which requires large η

Linear Regression

Let $W = (w_0, \dots, w_n)$, $X = (1, x_1, \dots, x_n)$,

$$l_W(X) = \sum w_i x_i$$

Given a training set $(X_0, y_0), \dots, (X_k, y_k)$

$$W_{\text{opt}} = \underset{W}{\operatorname{argmin}} \sum (y_j - l_W(X_j))^2 = \operatorname{argmin}_W \|y - M \cdot W\|^2$$

$$W_{\text{opt}} = (M^T \cdot M)^{-1} M^T y$$

where $M = [X_0, \dots, X_k]$

Linear Regression via Gradient Descent

Given a training set $(X, y) \in (X_0, y_0), \dots (X_k, y_k)$

$$J(W) = \frac{1}{2} (y - l(X))^2$$

$$\frac{\partial J(W)}{\partial w_i} = (y - l(X)) \cdot x_i$$


Widrow-Hoff learning rule:

$$w_i \leftarrow w_i - \eta \cdot \frac{\partial J(W)}{\partial w_i} = w_i - \eta \cdot (y - l(X)) \cdot x_i$$

Logistic Regression

Sigmoid function:

$$g(z) = \frac{1}{1 + e^{-z}}$$



$$h_W(X) = g(W \cdot X) = \frac{1}{1 + e^{-W \cdot X}}$$

Gradient Descent

Given a training set $(X_0, y_0), \dots (X_k, y_k)$

Assume:

$$P(y_j = 1 | X_j, W) = h_W(X_j)^{y_j} (1 - h_W(X_j)^{1-y_j})$$

$$L(W) = \prod_j P(y_j = 1 | X_j, W) = \prod_j h_W(X_j)^{y_j} (1 - h_W(X_j)^{1-y_j})$$

$$J(W) = \log L(W) = \sum_j y_j \log(h_W(X_j)) + (1 - y_j) \log(1 - h_W(X_j))$$

$$\frac{\partial J}{\partial w_i}(X) = (y - h_W(X_j))x_i \quad \text{using} \quad \frac{\partial g}{\partial z} = g(z)(1 - g(z))$$

Gradient Descent

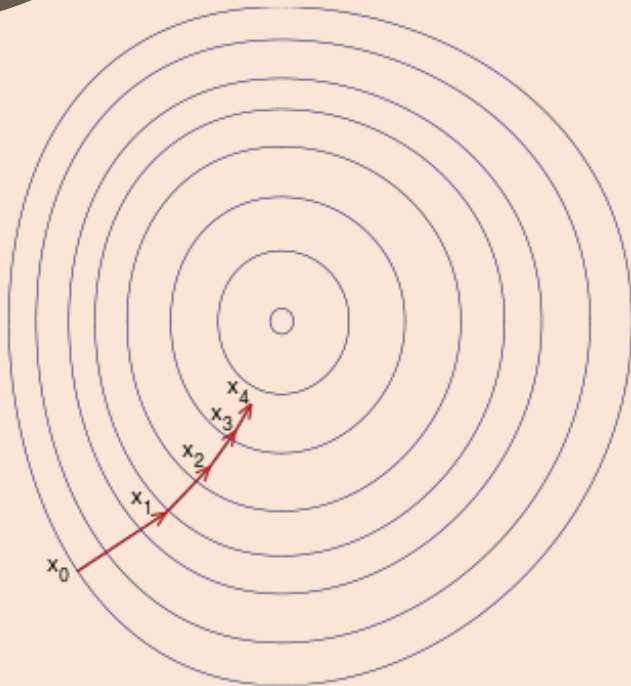
Given a training set $(X, y) \in (X_0, y_0), \dots (X_k, y_k)$

$$W \leftarrow (w_{0,1}, \dots w_{0,n})$$

for (X, y) *in* $(X_0, y_0), \dots (X_k, y_k)$:


$$h \leftarrow g(\sum w_i x_i)$$

$$w_i \leftarrow w_i + \eta \cdot x_i \cdot (y - h)$$



This shows the update algorithm used for the perceptron would converge to the weights maximizing likelihood

Feedforward Neural Networks



Deep Feedforward Neural Networks

